

Market Segmentation Analysis: Step 5 - Extracting Segments

Data-Driven Approaches to Consumer Clustering

November 8, 2025

Abstract

Data-driven market segmentation analysis is fundamentally exploratory. Consumer datasets typically lack clear structure, with preferences distributed across multidimensional space rather than forming distinct natural clusters. This comprehensive chapter explores how different extraction algorithms impose unique structural assumptions on segmented data, revealing that segmentation results depend equally on the underlying data characteristics and the chosen algorithmic approach. We examine distance-based methods (hierarchical clustering, k-means, self-organizing maps), model-based approaches (finite mixture models), and specialized algorithms with integrated variable selection. Through rigorous mathematical foundations, practical Python implementations, and detailed case studies using real consumer data, this work demonstrates that no single "best" algorithm exists—rather, each method offers distinct advantages for specific data characteristics and business objectives.

Contents

1	Introduction: The Power of Algorithmic Choice	4
1.1	Why Algorithm Choice Matters: The Spiral Example	4
1.2	Guiding Principles for Algorithm Selection	5
2	Grouping Consumers: Core Concepts	6
2.1	The Nature of Consumer Data	6
2.2	Distance Measures: Foundations of Similarity	6
2.2.1	Common Distance Measures	7
3	Distance-Based Methods	8
3.1	Overview of Distance-Based Approaches	8
3.2	Hierarchical Clustering Methods	8
3.2.1	Agglomerative vs Divisive Strategies	8
3.2.2	Linkage Methods: Defining Distance Between Groups	9
3.2.3	The Dendrogram: Visualizing Hierarchical Structure	11
3.2.4	Python Implementation: Hierarchical Clustering	11
3.3	Partitioning Methods	12
3.3.1	The k-Means Algorithm	12
3.3.2	Python Implementation: k-Means Clustering	14
3.4	Self-Organizing Maps (SOM)	16
3.4.1	SOM Architecture and Training	16
3.4.2	Python Implementation: Self-Organizing Maps	17

4 Model-Based Methods	20
4.1 Finite Mixture Models: Probabilistic Segmentation	20
4.2 The EM Algorithm	20
4.2.1 Model Selection Criteria	21
4.2.2 Python Implementation: Gaussian Mixture Models	22
5 Specialized Segmentation Methods	24
5.1 Methods with Integrated Variable Selection	24
5.1.1 Variable Selection for Binary Data (VSBD)	25
5.2 Bagged Clustering: Ensemble Approach	26
5.2.1 Python Implementation: Bagged Clustering	26
6 Determining the Optimal Number of Segments	29
6.1 The Elbow Method	29
6.2 Stability Analysis: The Gold Standard	29
6.2.1 Global Stability Analysis	29
6.2.2 Python Implementation: Stability Analysis	30
7 Practical Recommendations	32
8 Conclusion	33
9 Case Studies: From Theory to Practice	33
9.1 Case Study 1: Tourist Vacation Activities	33
9.1.1 Data Structure and Initial Exploration	34
9.1.2 Distance Matrix Analysis	34
9.1.3 Hierarchical Clustering Analysis	35
9.1.4 Segment Profiling	36
9.2 Case Study 2: Austrian Winter Vacation Activities	37
9.2.1 Dataset Overview	38
9.2.2 Bagged Clustering Implementation	39
9.2.3 Segment Profiles and Business Insights	40
10 Advanced Stability Analysis	40
10.1 Segment-Level Stability Within Solutions (SLSw)	40
10.1.1 SLSw Algorithm	41
10.1.2 Python Implementation: SLSw Analysis	42
10.2 Segment-Level Stability Across Solutions (SLSA)	44
10.2.1 SLSA with Entropy-Based Stability	45
10.3 Response Style Detection	45
10.3.1 Detecting Response Style Segments	45
10.3.2 Python Implementation: Response Style Detection	46
10.4 Handling Mixed Data Types	48
10.4.1 Gower Distance for Mixed Data	48
11 Complete Market Segmentation Workflow	51
11.1 End-to-End Process Overview	51
11.2 Decision Tree for Algorithm Selection	52
11.3 Stage-by-Stage Checklists	52
11.4 Common Pitfalls and Solutions	54

11.5 Segment Profiling Best Practices	54
11.5.1 Python Implementation: Automated Profiling Report	55
11.6 From Segmentation to Targeting Strategy	58
11.7 Reporting to Stakeholders	58
11.8 Software Implementation Guide	59
12 Final Recommendations	59
13 Conclusion: The Art and Science of Segmentation	60

1 Introduction: The Power of Algorithmic Choice

The Central Challenge

Market segmentation analysis faces a fundamental paradox: consumer data rarely contains naturally distinct groups, yet businesses require clear, actionable market segments. This chapter addresses the critical question:

How do different algorithms shape the segments they extract, and how can we select methods aligned with our business objectives?

The answer lies in understanding that:

- Segmentation methods **impose structure** on unstructured data
- Algorithm selection is as important as data quality
- Multiple valid segmentation solutions often exist
- Systematic exploration across methods is essential

1.1 Why Algorithm Choice Matters: The Spiral Example

Consider Figure 1, which demonstrates how dramatically different algorithms can produce different results from identical data.

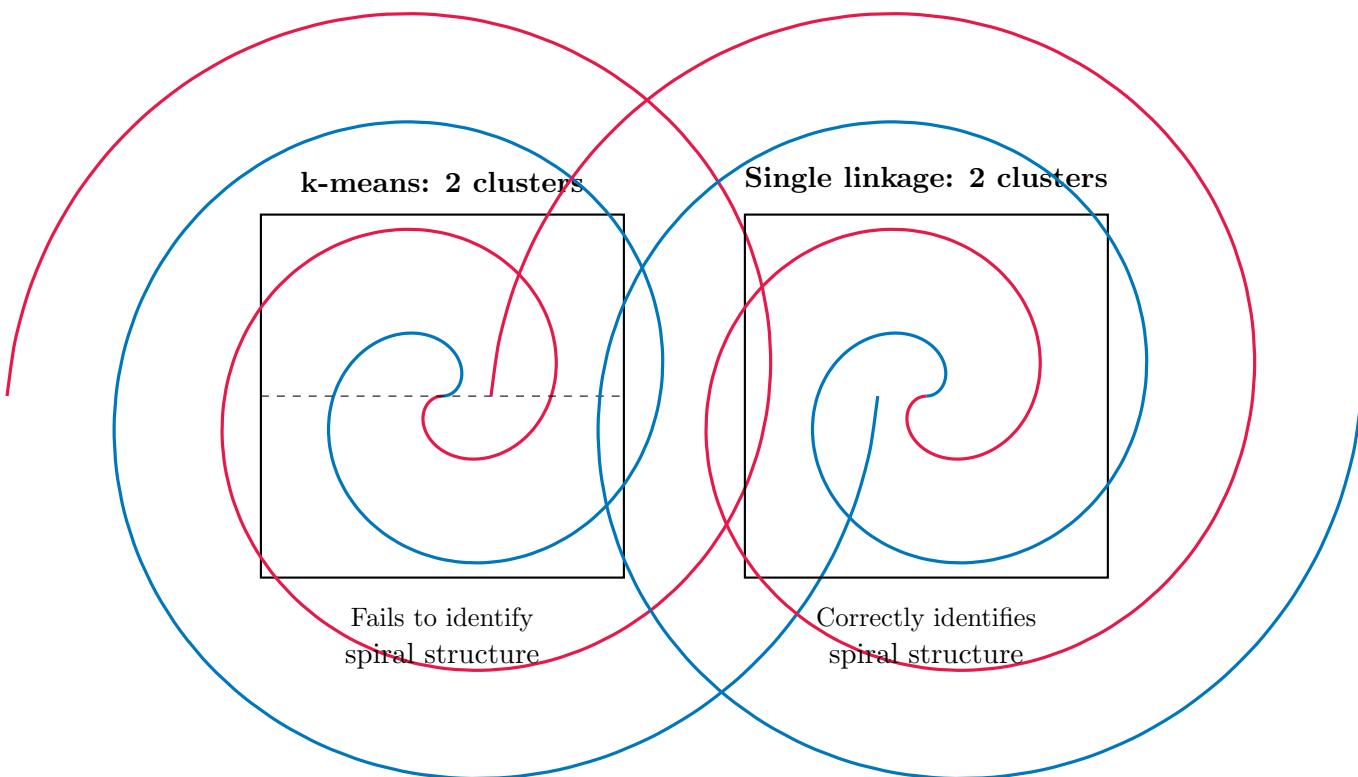


Figure 1: Comparison of k-means and single linkage clustering on spiral data. k-means imposes compact, spherical cluster assumptions and fails to detect the spiral structure. Single linkage allows snake-shaped clusters and correctly identifies both spirals.

Critical Insight

The spiral example illustrates a fundamental truth:

- **k-means** assumes compact, equally-sized, spherical clusters
- **Single linkage** allows elongated, chain-like structures
- Neither algorithm is universally "better"—each has optimal use cases
- The data structure determines which assumptions are appropriate

For consumer data, which typically lacks well-defined structure, algorithmic tendencies become even more influential in shaping the final segmentation solution.

1.2 Guiding Principles for Algorithm Selection

Algorithm Selection Framework

Three key dimensions inform extraction algorithm selection:

1. Data Characteristics

- Sample size: n consumers, p variables
- Variable scale: nominal, ordinal, metric, or mixed
- Data structure: presence of outliers, niche segments, hierarchical relationships
- Special features: temporal data, spatial data, repeated measurements

2. Desired Segment Properties

- Within-segment similarity: what should consumers share?
- Between-segment differences: how should segments differ?
- Segment size constraints: minimum viable segment size
- Number of segments: predetermined vs. exploratory

3. Computational Constraints

- Available memory for distance matrix computation
- Processing time requirements
- Software capabilities and available implementations

Table 1: Matching Data and Segment Characteristics to Algorithm Families

Scenario	Recommended Approach	Rationale
Small datasets ($n < 500$)	Hierarchical clustering	Dendrogram aids interpretation; memory not limiting
Large datasets ($n > 5000$)	Partitioning methods (k-means, SOM)	Computational efficiency; scalability
Mixed variable types	Model-based (finite mixtures)	Accommodates different distributions
Binary data with noise	Biclustering, VSBD	Simultaneous variable selection
Niche segment detection	Bagged clustering	Bootstrap resampling reveals rare segments
Unknown segment count	Hierarchical + dendrogram	Visual inspection suggests optimal k
Clear segment count	Partitioning methods	Direct optimization for specified k

2 Grouping Consumers: Core Concepts

2.1 The Nature of Consumer Data

Unstructured Consumer Preferences

Consumer data presents unique challenges:

Characteristics of Real Consumer Datasets:

1. **Continuous distributions:** Preferences spread across entire feature space
2. **No clear boundaries:** Gradual transitions rather than distinct clusters
3. **High dimensionality:** Many product attributes or behavioral variables
4. **Heterogeneity:** Consumers truly differ in individual ways

Implications for Segmentation:

- Algorithms must *create* structure from unstructured data
- Multiple valid segmentation solutions can coexist
- Domain knowledge must guide algorithm selection
- Validation through stability analysis is essential

2.2 Distance Measures: Foundations of Similarity

All distance-based segmentation methods require a mathematical definition of similarity or dissimilarity between consumers.

Distance Measure Requirements

A function $d : \mathcal{X} \times \mathcal{X} \rightarrow \mathbb{R}_{\geq 0}$ is a valid distance measure if it satisfies:

$$d(\mathbf{x}, \mathbf{y}) = d(\mathbf{y}, \mathbf{x}) \quad (\text{Symmetry}) \quad (1)$$

$$d(\mathbf{x}, \mathbf{y}) = 0 \Leftrightarrow \mathbf{x} = \mathbf{y} \quad (\text{Identity}) \quad (2)$$

$$d(\mathbf{x}, \mathbf{z}) \leq d(\mathbf{x}, \mathbf{y}) + d(\mathbf{y}, \mathbf{z}) \quad (\text{Triangle inequality}) \quad (3)$$

where $\mathbf{x}, \mathbf{y}, \mathbf{z} \in \mathbb{R}^p$ are p -dimensional consumer profiles.

2.2.1 Common Distance Measures

Let $\mathbf{x} = (x_1, \dots, x_p)$ and $\mathbf{y} = (y_1, \dots, y_p)$ represent two consumer profiles.

Table 2: Primary Distance Measures for Market Segmentation

Distance Type	Formula	Use Case
Euclidean	$d(\mathbf{x}, \mathbf{y}) = \sqrt{\sum_{j=1}^p (x_j - y_j)^2}$	Most common; metric data
Manhattan	$d(\mathbf{x}, \mathbf{y}) = \sum_{j=1}^p x_j - y_j $	Robust to outliers
Asymmetric Binary	$d(\mathbf{x}, \mathbf{y}) = \frac{\sum_{j:x_j=1 \vee y_j=1} \mathbb{1}_{x_j \neq y_j}}{\sum_{j:x_j=1 \vee y_j=1} 1}$	Activity patterns

Choosing Distance Measures

Euclidean Distance:

- Penalizes large differences more than small ones (squared terms)
- Assumes all dimensions equally important
- Requires standardization if variables have different scales
- Default choice for continuous consumer data

Manhattan Distance:

- Treats all differences equally (absolute values)
- More robust to outliers than Euclidean
- Natural for grid-like preference structures
- Often similar results to Euclidean on consumer data

Asymmetric Binary Distance:

- Critical for activity-based segmentation

- Shared presence (both = 1) indicates similarity
- Shared absence (both = 0) ignored as uninformative
- Example: Tourists who both engage in horseback riding are similar; tourists who both don't ride are not necessarily similar

3 Distance-Based Methods

3.1 Overview of Distance-Based Approaches

Core Principle

Distance-based methods identify market segments by:

1. Defining a mathematical measure of similarity (distance) between consumers
2. Finding groups where **within-segment similarity** is maximized
3. Ensuring **between-segment dissimilarity** is maximized

These methods translate the conceptual goal of "finding similar consumers" into precise mathematical optimization problems.

3.2 Hierarchical Clustering Methods

Hierarchical clustering mimics human intuition: we naturally think about grouping objects by progressively merging similar items or progressively dividing collections into subgroups.

3.2.1 Agglomerative vs Divisive Strategies

Table 3: Comparison of Hierarchical Clustering Strategies

Aspect	Agglomerative (Bottom-Up)	Divisive (Top-Down)
Starting point	Each consumer = own segment (n singletons)	All consumers in one segment
Process	Merge closest pairs iteratively	Split segments recursively
Ending point	One segment containing all consumers	Each consumer = own segment
Computational cost	$O(n^2 \log n)$	Typically higher
Most common	Yes (standard implementations)	No (rarely used)

Agglomerative hierarchical clustering is overwhelmingly preferred in practice because:

- Well-established, efficient algorithms exist (Lance & Williams framework)
- Deterministic: same data always produces same result
- Visualization through dendograms aids interpretation
- Standard implementations available in all software packages

3.2.2 Linkage Methods: Defining Distance Between Groups

The critical decision in hierarchical clustering: *how do we measure distance between two groups of consumers?*

Linkage Method Definitions

Let X and Y be two sets of consumer profiles. Define distance $d(X, Y)$ between these sets:

Single Linkage (Nearest Neighbor):

$$d_{\text{single}}(X, Y) = \min_{\mathbf{x} \in X, \mathbf{y} \in Y} d(\mathbf{x}, \mathbf{y}) \quad (4)$$

Complete Linkage (Furthest Neighbor):

$$d_{\text{complete}}(X, Y) = \max_{\mathbf{x} \in X, \mathbf{y} \in Y} d(\mathbf{x}, \mathbf{y}) \quad (5)$$

Average Linkage:

$$d_{\text{average}}(X, Y) = \frac{1}{|X| \cdot |Y|} \sum_{\mathbf{x} \in X} \sum_{\mathbf{y} \in Y} d(\mathbf{x}, \mathbf{y}) \quad (6)$$

Ward's Method (Minimum Variance):

$$d_{\text{Ward}}(X, Y) = \frac{|X| \cdot |Y|}{|X| + |Y|} \|\bar{\mathbf{x}} - \bar{\mathbf{y}}\|^2 \quad (7)$$

where $\bar{\mathbf{x}}$ and $\bar{\mathbf{y}}$ are cluster centroids (mean vectors).

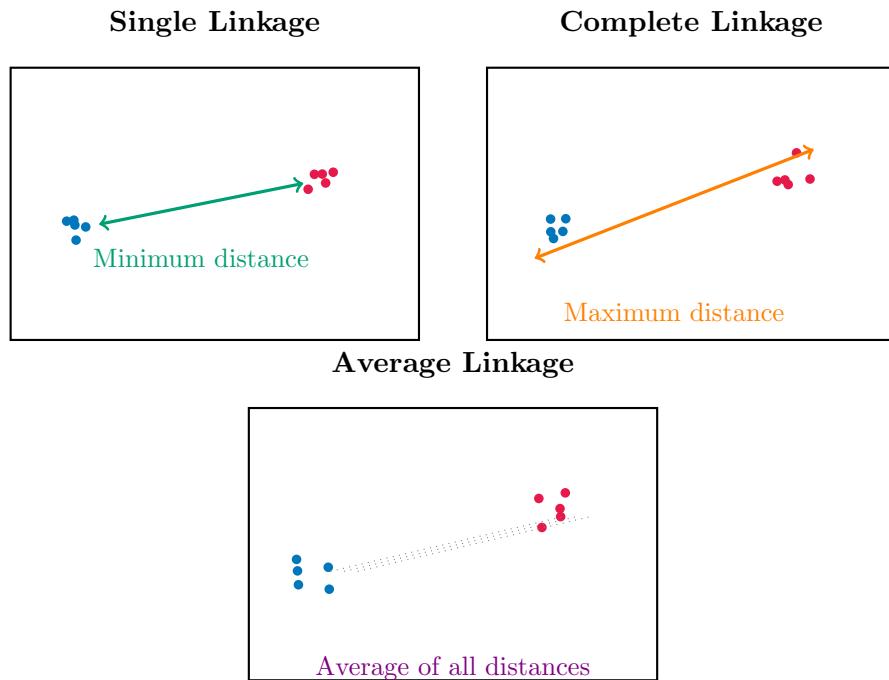


Figure 2: Visual comparison of linkage methods. Single linkage uses nearest neighbors, complete linkage uses furthest neighbors, and average linkage considers all pairwise distances.

Linkage Method Characteristics

Single Linkage:

- Detects elongated, chain-like, non-convex clusters (e.g., spirals)
- Sensitive to noise and outliers ("chaining" problem)
- Good for: Data with natural non-spherical structures

Complete Linkage:

- Produces compact, spherical clusters
- Less sensitive to outliers than single linkage
- Tends to create evenly-sized segments
- Good for: Well-separated consumer groups

Average Linkage:

- Compromise between single and complete
- Balances compactness with flexibility
- Often produces interpretable segments
- Good for: General-purpose consumer segmentation

Ward's Method:

- Minimizes within-cluster variance
- Strongly favors equal-sized, spherical clusters
- Most popular in market segmentation practice
- Uses squared Euclidean distance specifically
- Good for: Creating balanced segment portfolios

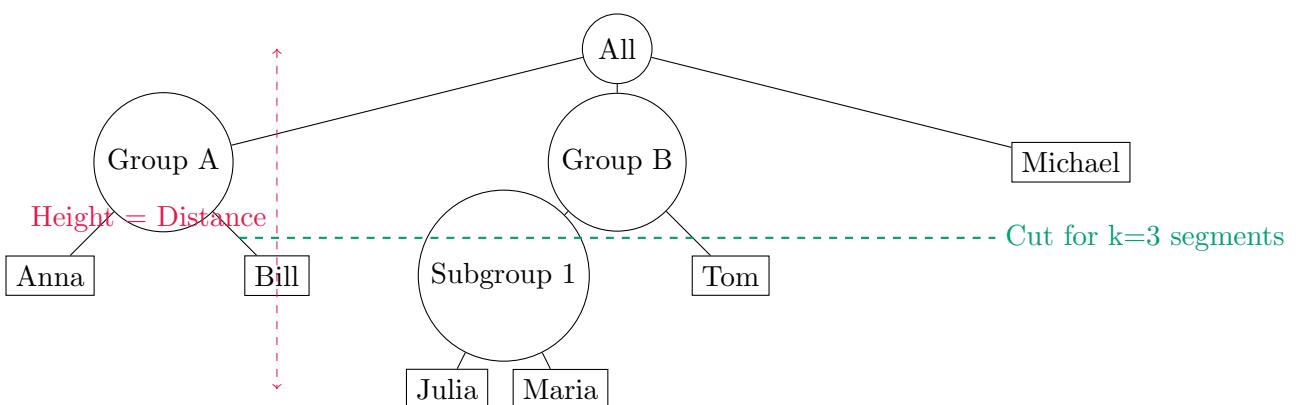
3.2.3 The Dendrogram: Visualizing Hierarchical Structure

Figure 3: Schematic dendrogram showing hierarchical cluster structure. Height represents distance between merged clusters. Horizontal line shows where to cut tree for desired number of segments.

Dendrogram Interpretation Challenges

For real consumer data, dendrograms rarely provide clear guidance on optimal segment count because:

- Consumer preferences are typically continuously distributed
- No natural "elbow" or clear height jump exists
- Leaf ordering is arbitrary (can permute left/right branches)
- Different software may display different-looking but equivalent dendrograms

Recommendation: Use dendrograms for rough exploration, but rely on systematic stability analysis (Section 7.5) for final segment count determination.

3.2.4 Python Implementation: Hierarchical Clustering

```

1 import numpy as np
2 import pandas as pd

```

```

3  from scipy.cluster.hierarchy import dendrogram, linkage, fcluster
4  from scipy.spatial.distance import pdist
5  import matplotlib.pyplot as plt
6  from sklearn.preprocessing import StandardScaler
7
8  # Example: Tourist vacation activities
9  data = pd.DataFrame({
10     'beach': [100, 100, 60, 70, 80, 0, 50],
11     'action': [0, 0, 40, 0, 0, 90, 20],
12     'culture': [0, 0, 0, 30, 20, 10, 30]
13 }, index=['Anna', 'Bill', 'Frank', 'Julia', 'Maria', 'Michael', 'Tom'])
14
15 # Standardize if needed (for variables on different scales)
16 # scaler = StandardScaler()
17 # data_scaled = scaler.fit_transform(data)
18
19 # Calculate pairwise distances (Manhattan distance)
20 distances = pdist(data, metric='cityblock') # 'cityblock' = Manhattan
21
22 # Perform hierarchical clustering (complete linkage)
23 linkage_matrix = linkage(distances, method='complete')
24
25 # Create dendrogram
26 plt.figure(figsize=(10, 6))
27 dendrogram(linkage_matrix, labels=data.index, leaf_font_size=12)
28 plt.xlabel('Tourist')
29 plt.ylabel('Distance')
30 plt.title('Complete Linkage Hierarchical Clustering')
31 plt.tight_layout()
32 plt.show()
33
34 # Extract k=3 segments
35 segments = fcluster(linkage_matrix, t=3, criterion='maxclust')
36 data['Segment'] = segments
37
38 # Display segment means
39 segment_means = data.groupby('Segment').mean()
40 print("\nSegment Characteristics:")
41 print(segment_means)

```

3.3 Partitioning Methods

Partitioning methods create a single division of consumers into k segments, optimizing directly for that specific number of segments rather than building a complete hierarchy.

3.3.1 The k-Means Algorithm

k-Means Clustering Algorithm

Input: Dataset $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$, number of segments k

Output: Partition S_1, \dots, S_k and centroids $\mathbf{c}_1, \dots, \mathbf{c}_k$

- 1: **Initialize:** Randomly select k observations as initial centroids $\mathbf{c}_1^{(0)}, \dots, \mathbf{c}_k^{(0)}$
- 2: Set iteration counter $t \leftarrow 0$

```

3: repeat
4:    $t \leftarrow t + 1$ 
5:   Assignment Step: Assign each  $\mathbf{x}_i$  to nearest centroid

```

$$S_j^{(t)} = \{\mathbf{x}_i : \|\mathbf{x}_i - \mathbf{c}_j^{(t-1)}\|^2 \leq \|\mathbf{x}_i - \mathbf{c}_h^{(t-1)}\|^2, \forall h \in \{1, \dots, k\}\}$$

```
6:   Update Step: Recompute centroids as segment means
```

$$\mathbf{c}_j^{(t)} = \frac{1}{|S_j^{(t)}|} \sum_{\mathbf{x}_i \in S_j^{(t)}} \mathbf{x}_i$$

```
7: until Centroids no longer change:  $\mathbf{c}_j^{(t)} = \mathbf{c}_j^{(t-1)}$  for all  $j$ 
```

```
8: return Segments  $S_1^{(t)}, \dots, S_k^{(t)}$  and centroids  $\mathbf{c}_1^{(t)}, \dots, \mathbf{c}_k^{(t)}$ 
```

Step 1: Random Init Step 2: Assignment

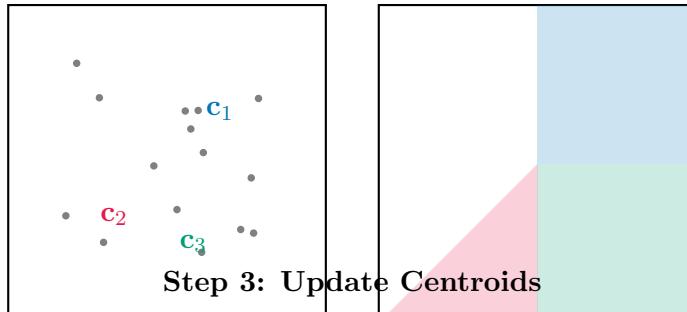


Figure 4: k-Means algorithm visualization: (1) Random initialization, (2) Assignment to nearest centroid, (3) Centroid update to segment means. Process repeats until convergence.

Key Properties of k-Means

Strengths:

- Computationally efficient: $O(nkp \cdot \text{iterations})$
- Scales to large datasets ($n > 10,000$)
- Simple to understand and implement

- Produces compact, spherical segments

Limitations:

- Requires pre-specification of k
- Sensitive to initialization (local optima)
- Assumes spherical clusters of similar size
- Only works with Euclidean distance
- Sensitive to outliers

Best Practices:

1. Run multiple times (10-50) with different random initializations
2. Select solution with minimum within-cluster sum of squares
3. Standardize variables if on different scales
4. Consider k-means++ initialization for better starting points

3.3.2 Python Implementation: k-Means Clustering

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.cluster import KMeans
4 from sklearn.preprocessing import StandardScaler
5 import matplotlib.pyplot as plt
6 from sklearn.metrics import silhouette_score
7
8 # Generate or load data
9 np.random.seed(42)
10 n_consumers = 500
11
12 # Artificial mobile phone data
13 features = np.concatenate([
14     np.random.normal(2, 0.5, 200),
15     np.random.normal(5, 0.5, 200),
16     np.random.normal(8, 0.5, 100)
17 ])
18 prices = np.concatenate([
19     np.random.normal(2, 0.5, 200),
20     np.random.normal(5, 0.5, 200),
21     np.random.normal(8, 0.5, 100)
22 ])
23
24 data = pd.DataFrame({'features': features, 'price': prices})
25
26 # Standardize (important for k-means!)
27 scaler = StandardScaler()
28 data_scaled = scaler.fit_transform(data)
```

```

29
30 # Elbow method: try k=1 to k=10
31 inertias = []
32 silhouette_scores = []
33 K_range = range(1, 11)
34
35 for k in K_range:
36     kmeans = KMeans(n_clusters=k, n_init=10, random_state=42)
37     kmeans.fit(data_scaled)
38     inertias.append(kmeans.inertia_)
39     if k > 1:
40         silhouette_scores.append(silhouette_score(data_scaled, kmeans.labels_))
41
42 # Plot elbow curve
43 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
44
45 ax1.plot(K_range, inertias, 'bo-')
46 ax1.set_xlabel('Number of Segments (k)')
47 ax1.set_ylabel('Within-Cluster Sum of Squares')
48 ax1.set_title('Elbow Method')
49 ax1.grid(True)
50
51 ax2.plot(range(2, 11), silhouette_scores, 'ro-')
52 ax2.set_xlabel('Number of Segments (k)')
53 ax2.set_ylabel('Silhouette Score')
54 ax2.set_title('Silhouette Analysis')
55 ax2.grid(True)
56
57 plt.tight_layout()
58 plt.show()
59
60 # Fit final model with k=3
61 kmeans_final = KMeans(n_clusters=3, n_init=50, random_state=42)
62 data['Segment'] = kmeans_final.fit_predict(data_scaled)
63
64 # Visualize segments
65 plt.figure(figsize=(8, 6))
66 for segment in range(3):
67     segment_data = data[data['Segment'] == segment]
68     plt.scatter(segment_data['features'], segment_data['price'],
69                 label=f'Segment {segment+1}', alpha=0.6, s=50)
70
71 # Plot centroids (transform back to original scale)
72 centroids_original = scaler.inverse_transform(kmeans_final.cluster_centers_)
73 plt.scatter(centroids_original[:, 0], centroids_original[:, 1],
74             c='black', marker='*', s=300, edgecolors='white', linewidth=2,
75             label='Centroids')
76
77 plt.xlabel('Features')
78 plt.ylabel('Price')
79 plt.title('k-Means Segmentation (k=3)')
80 plt.legend()
81 plt.grid(True, alpha=0.3)
82 plt.show()
83
84 # Segment profiles

```

```

85 print("\nSegment Profiles:")
86 print(data.groupby('Segment')[['features', 'price']].mean())
87 print("\nSegment Sizes:")
88 print(data['Segment'].value_counts().sort_index())

```

3.4 Self-Organizing Maps (SOM)

Self-Organizing Maps represent a fundamentally different approach to segmentation, borrowing from neural network methodology to create topology-preserving mappings of high-dimensional consumer data.

The SOM Principle

Self-Organizing Maps (Kohonen networks) simultaneously achieve two objectives:

1. **Dimensionality reduction:** Map high-dimensional consumer data onto a 2D grid
2. **Clustering:** Group similar consumers while preserving topological relationships

The key innovation: **neighboring segments on the map are more similar than distant segments**, providing intuitive visualization of market landscape.

3.4.1 SOM Architecture and Training

Input Layer

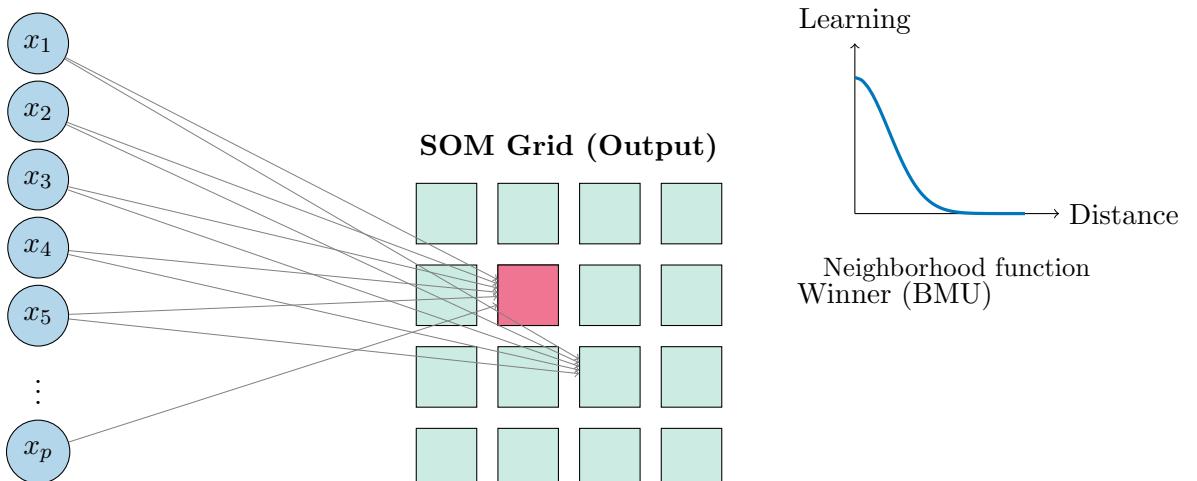


Figure 5: Self-Organizing Map architecture showing input layer (consumer features), 2D output grid, winner node (BMU - Best Matching Unit), and neighborhood function for competitive learning.

SOM Training Algorithm

Input: Dataset $X = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$, grid dimensions $m \times n$

Output: Trained weight vectors \mathbf{w}_{ij} for each grid node (i, j)

- 1: **Initialize:** Randomly initialize weight vectors $\mathbf{w}_{ij}^{(0)} \in \mathbb{R}^p$ for all grid nodes
- 2: Set learning rate $\alpha(0)$ and neighborhood radius $\sigma(0)$
- 3: **for** epoch $t = 1$ to T_{\max} **do**
- 4: **for** each consumer $\mathbf{x}_k \in X$ **do**
- 5: **Competition:** Find Best Matching Unit (BMU)

$$\text{BMU} = \arg \min_{i,j} \|\mathbf{x}_k - \mathbf{w}_{ij}^{(t)}\|$$

- 6: **Cooperation:** Compute neighborhood function

$$h_{ij,\text{BMU}}(t) = \exp\left(-\frac{\|\mathbf{r}_{ij} - \mathbf{r}_{\text{BMU}}\|^2}{2\sigma(t)^2}\right)$$

where \mathbf{r}_{ij} is grid position of node (i, j)

- 7: **Adaptation:** Update weights

$$\mathbf{w}_{ij}^{(t+1)} = \mathbf{w}_{ij}^{(t)} + \alpha(t) \cdot h_{ij,\text{BMU}}(t) \cdot (\mathbf{x}_k - \mathbf{w}_{ij}^{(t)})$$

- 8: **end for**

9: Decrease learning rate: $\alpha(t+1) = \alpha(t) \cdot \text{decay factor}$

10: Decrease neighborhood: $\sigma(t+1) = \sigma(t) \cdot \text{decay factor}$

- 11: **end for**

12: **return** Weight vectors \mathbf{w}_{ij}

SOM Advantages for Market Segmentation

Unique Benefits:

- **Visualization:** 2D map shows relationships between segments
- **Topology preservation:** Similar segments are neighbors on grid
- **Outlier detection:** Isolated nodes reveal niche segments
- **No predetermined k :** Segments emerge naturally from grid structure
- **Hierarchical view:** Can perform secondary clustering on SOM nodes

Practical Applications:

1. Initial exploratory analysis of consumer landscape
2. Identifying transition zones between major segments
3. Detecting emerging micro-segments
4. Creating intuitive management presentations

3.4.2 Python Implementation: Self-Organizing Maps

```
1 import numpy as np
2 import pandas as pd
3 from sklearn.preprocessing import StandardScaler
```

```

4 import matplotlib.pyplot as plt
5 from matplotlib.patches import Rectangle
6
7 # Note: Need to install minisom library
8 # pip install minisom
9 from minisom import MiniSom
10
11 # Generate sample consumer data
12 np.random.seed(42)
13 n_consumers = 500
14
15 # Three natural segments
16 seg1 = np.random.normal([2, 8], [0.5, 0.5], (200, 2))
17 seg2 = np.random.normal([8, 2], [0.5, 0.5], (200, 2))
18 seg3 = np.random.normal([5, 5], [0.7, 0.7], (100, 2))
19
20 data = np.vstack([seg1, seg2, seg3])
21 data_df = pd.DataFrame(data, columns=['feature_preference',
22                                     'price_sensitivity'])
23
24 # Standardize data
25 scaler = StandardScaler()
26 data_scaled = scaler.fit_transform(data)
27
28 # Initialize and train SOM
29 som_shape = (10, 10) # 10x10 grid
30 som = MiniSom(som_shape[0], som_shape[1], data_scaled.shape[1],
31                 sigma=1.5, # Initial neighborhood radius
32                 learning_rate=0.5,
33                 neighborhood_function='gaussian')
34
35 # Initialize weights
36 som.random_weights_init(data_scaled)
37
38 # Train SOM
39 print("Training SOM...")
40 som.train_random(data_scaled, num_iteration=1000)
41 print("Training complete!")
42
43 # Visualize SOM
44 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 6))
45
46 # Distance map (U-matrix)
47 from matplotlib.collections import PatchCollection
48
49 distance_map = som.distance_map()
50 ax1.imshow(distance_map.T, cmap='bone_r', origin='lower')
51 ax1.set_title('U-Matrix (Unified Distance Matrix)', fontsize=14)
52 ax1.set_xlabel('SOM Grid X')
53 ax1.set_ylabel('SOM Grid Y')
54
55 # Add colorbar
56 cbar1 = plt.colorbar(ax1.images[0], ax=ax1)
57 cbar1.set_label('Average Distance to Neighbors')
58
59 # Hit map (frequency)

```

```

59 frequencies = np.zeros(som_shape)
60 for x in data_scaled:
61     winner = som.winner(x)
62     frequencies[winner] += 1
63
64 im2 = ax2.imshow(frequencies.T, cmap='YlOrRd', origin='lower')
65 ax2.set_title('Hit Map (Consumer Distribution)', fontsize=14)
66 ax2.set_xlabel('SOM Grid X')
67 ax2.set_ylabel('SOM Grid Y')
68
69 # Annotate hit counts
70 for i in range(som_shape[0]):
71     for j in range(som_shape[1]):
72         if frequencies[i, j] > 0:
73             ax2.text(i, j, int(frequencies[i, j]),
74                      ha='center', va='center', fontsize=8)
75
76 cbar2 = plt.colorbar(im2, ax=ax2)
77 cbar2.set_label('Number of Consumers')
78
79 plt.tight_layout()
80 plt.show()
81
82 # Assign consumers to SOM nodes
83 som_clusters = np.array([som.winner(x) for x in data_scaled])
84 data_df['som_x'] = som_clusters[:, 0]
85 data_df['som_y'] = som_clusters[:, 1]
86
87 # Perform secondary k-means clustering on SOM nodes
88 from sklearn.cluster import KMeans
89
90 # Get unique SOM node positions
91 unique_nodes = np.unique(som_clusters, axis=0)
92 node_weights = np.array([som.get_weights()[i, j]
93                         for i, j in unique_nodes])
94
95 # Cluster SOM nodes into k=3 macro-segments
96 kmeans_som = KMeans(n_clusters=3, random_state=42)
97 node_labels = kmeans_som.fit_predict(node_weights)
98
99 # Map back to consumers
100 node_to_label = {tuple(node): label for node, label
101                  in zip(unique_nodes, node_labels)}
102 data_df['segment'] = [node_to_label[tuple(pos)]
103                      for pos in som_clusters]
104
105 # Visualize final segmentation
106 fig, ax = plt.subplots(figsize=(10, 8))
107
108 colors = ['blue', 'red', 'green']
109 for segment in range(3):
110     segment_data = data_df[data_df['segment'] == segment]
111     ax.scatter(segment_data['feature_preference'],
112                segment_data['price_sensitivity'],
113                c=colors[segment], label=f'Segment {segment+1}',
114                alpha=0.6, s=50)

```

```

115
116 # Plot segment centroids
117 centroids_original = scaler.inverse_transform(kmeans_som.cluster_centers_)
118 ax.scatter(centroids_original[:, 0], centroids_original[:, 1],
119             c='black', marker='*', s=500, edgecolors='white',
120             linewidth=2, label='Centroids', zorder=5)
121
122 ax.set_xlabel('Feature Preference')
123 ax.set_ylabel('Price Sensitivity')
124 ax.set_title('SOM-based Market Segmentation')
125 ax.legend()
126 ax.grid(True, alpha=0.3)
127 plt.show()
128
129 # Segment profiles
130 print("\nSegment Profiles:")
131 print(data_df.groupby('segment')[['feature_preference',
132                                     'price_sensitivity']].mean())
133 print("\nSegment Sizes:")
134 print(data_df['segment'].value_counts().sort_index())

```

4 Model-Based Methods

4.1 Finite Mixture Models: Probabilistic Segmentation

Model-based segmentation represents a paradigm shift from distance-based approaches. Instead of measuring similarity through distances, these methods assume consumers are drawn from a mixture of probability distributions.

4.2 The EM Algorithm

Expectation-Maximization (EM) Algorithm

Purpose: Maximum likelihood estimation for finite mixture models

- 1: **Initialize:** Set initial parameter values $\Theta^{(0)}$
- 2: Set iteration counter $t \leftarrow 0$
- 3: **repeat**
- 4: $t \leftarrow t + 1$
- 5: **E-Step (Expectation):** Compute posterior probabilities

$$\tau_{ij}^{(t)} = P(z_i = j | \mathbf{y}_i, \Theta^{(t-1)})$$

for all consumers $i = 1, \dots, n$ and segments $j = 1, \dots, k$

- 6: **M-Step (Maximization):** Update parameters
- 7: Update segment sizes:

$$\pi_j^{(t)} = \frac{1}{n} \sum_{i=1}^n \tau_{ij}^{(t)}$$

8: Update segment-specific parameters $\boldsymbol{\theta}_j^{(t)}$ by maximizing:

$$\boldsymbol{\theta}_j^{(t)} = \arg \max_{\boldsymbol{\theta}_j} \sum_{i=1}^n \tau_{ij}^{(t)} \log f_j(\mathbf{y}_i | \boldsymbol{\theta}_j)$$

9: **until** Convergence: $|\ell(\boldsymbol{\Theta}^{(t)}) - \ell(\boldsymbol{\Theta}^{(t-1)})| < \epsilon$

10: **return** Parameters $\boldsymbol{\Theta}^{(t)}$

where $\ell(\boldsymbol{\Theta})$ is the log-likelihood function.

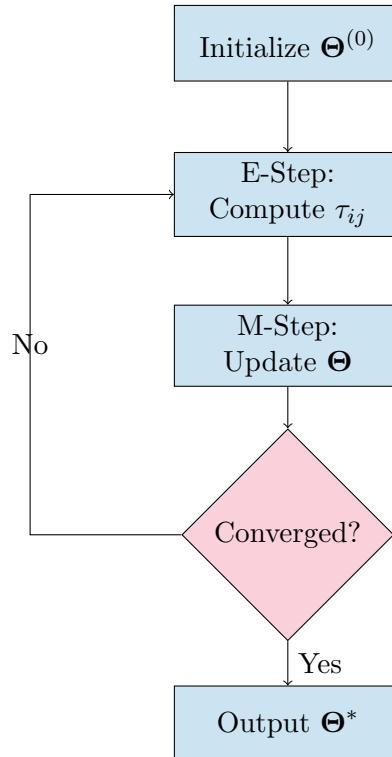


Figure 6: Flowchart of the Expectation-Maximization (EM) algorithm for estimating finite mixture model parameters. The algorithm iterates between computing membership probabilities (E-step) and updating parameters (M-step) until convergence.

4.2.1 Model Selection Criteria

Information Criteria for Model Selection

When number of segments k is unknown, compare models using:

Akaike Information Criterion (AIC):

$$AIC = -2\ell(\boldsymbol{\Theta}) + 2d \quad (10)$$

Bayesian Information Criterion (BIC):

$$BIC = -2\ell(\boldsymbol{\Theta}) + d \log(n) \quad (11)$$

Integrated Completed Likelihood (ICL):

$$\text{ICL} = \text{BIC} - 2 \sum_{i=1}^n \sum_{j=1}^k \tau_{ij} \log(\tau_{ij}) \quad (12)$$

where:

- $\ell(\Theta)$: Maximized log-likelihood
- d : Number of free parameters
- n : Sample size
- τ_{ij} : Posterior probability (from Eq. ??)

Decision rule: Select model minimizing criterion value

Comparison:

- BIC penalizes complexity more heavily than AIC ($\log(n) > 2$ for $n > 7$)
- ICL additionally penalizes unclear segment assignments (entropy term)
- No single criterion universally superior—compare multiple

4.2.2 Python Implementation: Gaussian Mixture Models

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.mixture import GaussianMixture
4 from sklearn.preprocessing import StandardScaler
5 import matplotlib.pyplot as plt
6 from scipy import stats
7
8 # Generate sample data
9 np.random.seed(42)
10
11 # Three segments with different means and covariances
12 seg1 = np.random.multivariate_normal([2, 2], [[1, 0.5], [0.5, 1]], 200)
13 seg2 = np.random.multivariate_normal([8, 8], [[1, -0.3], [-0.3, 1]], 200)
14 seg3 = np.random.multivariate_normal([5, 8], [[2, 0], [0, 0.5]], 100)
15
16 data = np.vstack([seg1, seg2, seg3])
17 data_df = pd.DataFrame(data, columns=['features', 'price'])
18
19 # Standardize
20 scaler = StandardScaler()
21 data_scaled = scaler.fit_transform(data)
22
23 # Model selection: try k=1 to k=8
24 k_range = range(1, 9)
25 bic_scores = []
26 aic_scores = []
27

```

```

28 for k in k_range:
29     gmm = GaussianMixture(n_components=k, covariance_type='full',
30                           n_init=10, random_state=42)
31     gmm.fit(data_scaled)
32     bic_scores.append(gmm.bic(data_scaled))
33     aic_scores.append(gmm.aic(data_scaled))
34
35 # Plot information criteria
36 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(12, 4))
37
38 ax1.plot(k_range, bic_scores, 'bo-', linewidth=2, markersize=8)
39 ax1.set_xlabel('Number of Segments (k)')
40 ax1.set_ylabel('BIC')
41 ax1.set_title('Bayesian Information Criterion')
42 ax1.grid(True, alpha=0.3)
43 ax1.axvline(x=np.argmin(bic_scores)+1, color='red',
44              linestyle='--', label='Optimal k')
45 ax1.legend()
46
47 ax2.plot(k_range, aic_scores, 'ro-', linewidth=2, markersize=8)
48 ax2.set_xlabel('Number of Segments (k)')
49 ax2.set_ylabel('AIC')
50 ax2.set_title('Akaike Information Criterion')
51 ax2.grid(True, alpha=0.3)
52 ax2.axvline(x=np.argmin(aic_scores)+1, color='blue',
53              linestyle='--', label='Optimal k')
54 ax2.legend()
55
56 plt.tight_layout()
57 plt.show()
58
59 # Fit optimal model (k=3)
60 gmm_optimal = GaussianMixture(n_components=3, covariance_type='full',
61                               n_init=50, random_state=42)
62 data_df['segment'] = gmm_optimal.fit_predict(data_scaled)
63 data_df['max_probability'] = np.max(gmm_optimal.predict_proba(data_scaled),
64                                   axis=1)
65
66 # Visualize with uncertainty
67 fig, ax = plt.subplots(figsize=(10, 8))
68
69 colors = ['blue', 'red', 'green']
70 for segment in range(3):
71     segment_data = data_df[data_df['segment'] == segment]
72
73     # Size reflects uncertainty (larger = more uncertain)
74     sizes = 100 * (1 - segment_data['max_probability'])
75
76     ax.scatter(segment_data['features'], segment_data['price'],
77                c=colors[segment], s=sizes, alpha=0.6,
78                label=f'Segment {segment+1}', edgecolors='black', linewidth=0.5)
79
80 # Plot Gaussian ellipses
81 from matplotlib.patches import Ellipse
82 for i in range(3):

```

```

83     mean = scaler.inverse_transform(gmm_optimal.means_[i:i+1])[0]
84     covariance = gmm_optimal.covariances_[i]
85
86     # Transform covariance back to original scale
87     # This is approximate for visualization
88     v, w = np.linalg.eigh(covariance)
89     angle = np.arctan2(w[1, 0], w[0, 0]) * 180 / np.pi
90
91     # 2 standard deviations
92     for n_std in [1, 2]:
93         ell = Ellipse(mean,
94                         width=2*n_std*np.sqrt(v[0]),
95                         height=2*n_std*np.sqrt(v[1]),
96                         angle=angle,
97                         edgecolor=colors[i],
98                         facecolor='none',
99                         linewidth=2,
100                        linestyle='--')
101        ax.add_patch(ell)
102
103 ax.set_xlabel('Features')
104 ax.set_ylabel('Price')
105 ax.set_title('Gaussian Mixture Model Segmentation\n(Point size indicates
106    ↪ uncertainty)')
107 ax.legend()
108 ax.grid(True, alpha=0.3)
109 plt.show()
110
111 # Segment profiles
112 print("\nSegment Profiles (Original Scale):")
113 print(data_df.groupby('segment')[['features', 'price']].agg(['mean', 'std']))
114 print("\nSegment Sizes:")
115 print(data_df['segment'].value_counts().sort_index())
116 print("\nAverage Posterior Probability by Segment:")
117 print(data_df.groupby('segment')['max_probability'].mean())

```

5 Specialized Segmentation Methods

5.1 Methods with Integrated Variable Selection

Standard clustering algorithms use all provided segmentation variables. For high-dimensional consumer data, however, many variables may be irrelevant or noisy, obscuring true market segments.

The Variable Selection Challenge

Consumer datasets often contain:

- **Relevant variables** that define true market segments
- **Irrelevant variables** that add noise
- **Redundant variables** that duplicate information

Including irrelevant variables can:

1. Mask true segment structure
2. Reduce statistical power
3. Increase computational cost
4. Decrease interpretability

5.1.1 Variable Selection for Binary Data (VSBD)

VSBD Optimization Problem

For binary segmentation variables, VSBD simultaneously:

1. Selects subset of $V \subseteq \{1, \dots, p\}$ relevant variables
2. Partitions consumers into k segments based only on V

Objective function: Maximize Ratkowsky-Lance index

$$RL(V, k) = \frac{1}{\sqrt{k}} \cdot \frac{1}{|V|} \sum_{j \in V} \frac{B_j}{T_j} \quad (13)$$

where:

- $B_j = \sum_{h=1}^k n_h (\bar{y}_{hj} - \bar{y}_j)^2$: Between-segment variation for variable j
- $T_j = \sum_{i=1}^n (y_{ij} - \bar{y}_j)^2$: Total variation for variable j
- n_h : Size of segment h

VSBD Advantages

Key Benefits:

- Identifies **niche segments** defined by specific activity combinations
- Removes noisy variables automatically
- Increases interpretability (fewer variables per segment)
- Particularly effective for binary activity data

Typical Applications:

- Vacation activity segmentation
- Product feature preferences
- Media consumption patterns
- Purchase behavior clustering

5.2 Bagged Clustering: Ensemble Approach

Bagged clustering combines bootstrap aggregation with hierarchical and partitioning methods to improve stability and niche segment detection.

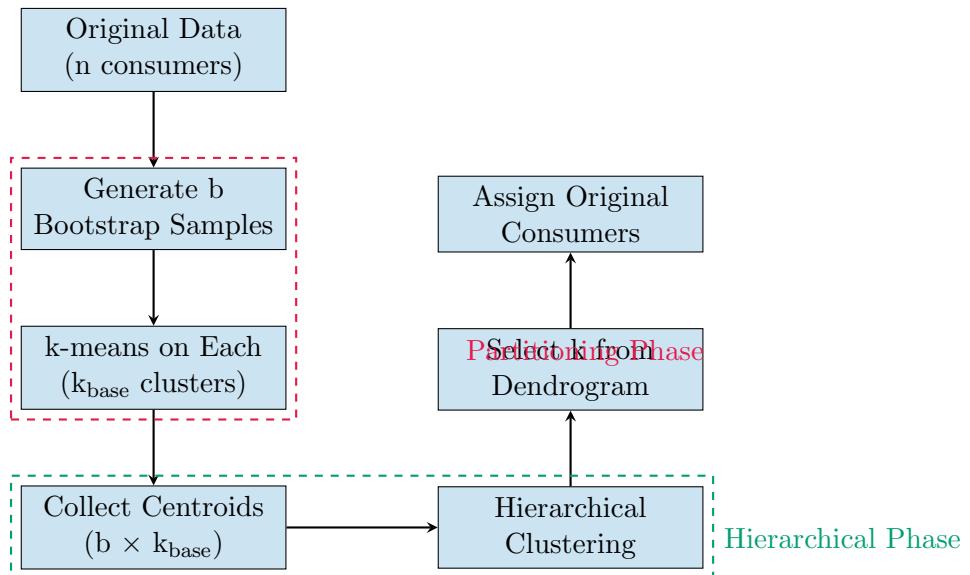
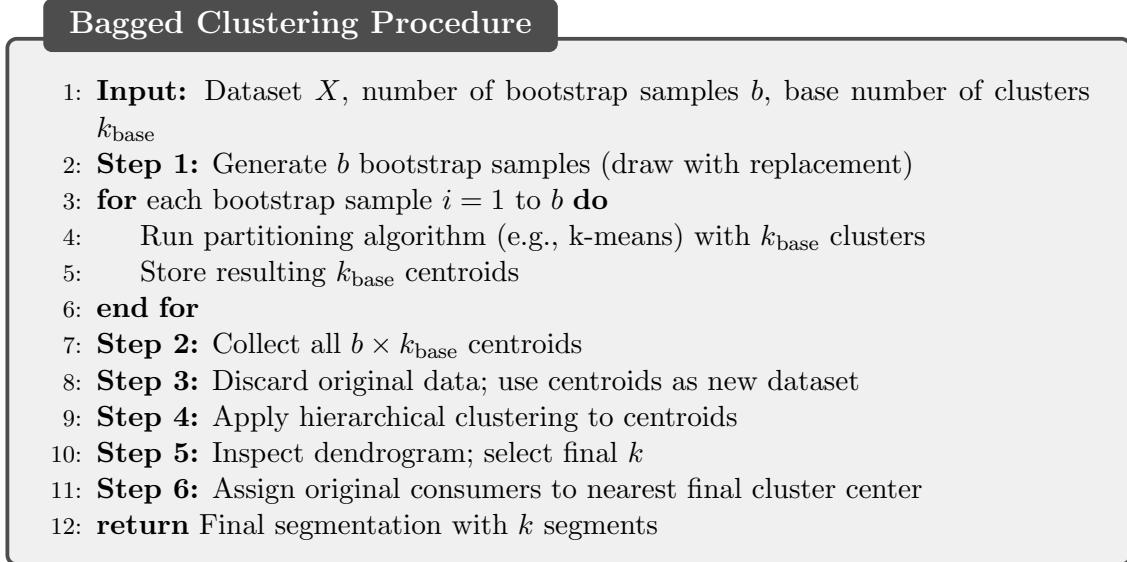


Figure 7: Bagged clustering workflow combining bootstrap resampling, partitioning methods, and hierarchical clustering to improve niche segment detection.

5.2.1 Python Implementation: Bagged Clustering

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.cluster import KMeans
4 from scipy.cluster.hierarchy import dendrogram, linkage, fcluster
5 from sklearn.preprocessing import StandardScaler
6 import matplotlib.pyplot as plt
7

```

```

8 # Generate sample data
9 np.random.seed(42)
10 n_consumers = 800
11
12 # Main segments + niche segment
13 main1 = np.random.normal([2, 2], [1, 1], (350, 2))
14 main2 = np.random.normal([8, 8], [1, 1], (350, 2))
15 niche = np.random.normal([5, 1], [0.3, 0.3], (100, 2)) # Small niche
16
17 data = np.vstack([main1, main2, niche])
18 true_labels = np.array([0]*350 + [1]*350 + [2]*100)
19
20 # Standardize
21 scaler = StandardScaler()
22 data_scaled = scaler.fit_transform(data)
23
24 # Bagged clustering parameters
25 n_bootstrap = 50
26 base_k = 10
27
28 # Store all centroids from bootstrap samples
29 all_centroids = []
30
31 print("Running bagged clustering...")
32 for b in range(n_bootstrap):
33     # Bootstrap sample
34     indices = np.random.choice(len(data_scaled), size=len(data_scaled),
35                                 replace=True)
36     bootstrap_data = data_scaled[indices]
37
38     # k-means with base_k clusters
39     kmeans = KMeans(n_clusters=base_k, n_init=10, random_state=b)
40     kmeans.fit(bootstrap_data)
41
42     # Store centroids
43     all_centroids.append(kmeans.cluster_centers_)
44
45 # Stack all centroids
46 all_centroids = np.vstack(all_centroids)
47 print(f"Collected {len(all_centroids)} centroids")
48
49 # Hierarchical clustering on centroids
50 linkage_matrix = linkage(all_centroids, method='average')
51
52 # Plot dendrogram
53 plt.figure(figsize=(12, 6))
54 dendrogram(linkage_matrix)
55 plt.xlabel('Centroid Index')
56 plt.ylabel('Distance')
57 plt.title('Bagged Clustering Dendrogram')
58 plt.axhline(y=2.5, color='r', linestyle='--', label='Cut line for k=3')
59 plt.legend()
60 plt.tight_layout()
61 plt.show()
62
63 # Extract k=3 segments from dendrogram

```

```

64 final_k = 3
65 centroid_labels = fcluster(linkage_matrix, t=final_k, criterion='maxclust')
66
67 # For each centroid, compute its final segment
68 # Then assign original consumers to segments
69 from scipy.spatial.distance import cdist
70
71 # Get final segment centers (mean of centroids in each segment)
72 final_centers = np.array([all_centroids[centroid_labels == j].mean(axis=0)
73                         for j in range(1, final_k+1)])
74
75 # Assign original consumers to nearest final center
76 distances = cdist(data_scaled, final_centers)
77 consumer_segments = np.argmin(distances, axis=1)
78
79 # Visualize results
80 fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 6))
81
82 # Original data with true labels
83 scatter1 = ax1.scatter(data[:, 0], data[:, 1], c=true_labels,
84                        cmap='viridis', alpha=0.6, s=50)
85 ax1.set_title('True Segments')
86 ax1.set_xlabel('Feature 1')
87 ax1.set_ylabel('Feature 2')
88 plt.colorbar(scatter1, ax=ax1)
89
90 # Bagged clustering results
91 scatter2 = ax2.scatter(data[:, 0], data[:, 1], c=consumer_segments,
92                        cmap='viridis', alpha=0.6, s=50)
93 final_centers_orig = scaler.inverse_transform(final_centers)
94 ax2.scatter(final_centers_orig[:, 0], final_centers_orig[:, 1],
95             c='red', marker='*', s=500, edgecolors='white', linewidth=2,
96             label='Final Centers')
97 ax2.set_title('Bagged Clustering Results')
98 ax2.set_xlabel('Feature 1')
99 ax2.set_ylabel('Feature 2')
100 ax2.legend()
101 plt.colorbar(scatter2, ax=ax2)
102
103 plt.tight_layout()
104 plt.show()
105
106 # Segment sizes
107 print("\nSegment Sizes:")
108 unique, counts = np.unique(consumer_segments, return_counts=True)
109 for seg, count in zip(unique, counts):
110     print(f"Segment {seg+1}: {count} consumers ({100*count/len(data):.1f}%)")

```

6 Determining the Optimal Number of Segments

6.1 The Elbow Method

Scree Plot Analysis

The **elbow method** plots within-cluster sum of squares (WCSS) against number of clusters k :

$$\text{WCSS}(k) = \sum_{j=1}^k \sum_{\mathbf{x}_i \in S_j} \|\mathbf{x}_i - \mathbf{c}_j\|^2 \quad (14)$$

Interpretation:

- WCSS always decreases with increasing k
- Look for "elbow" where decrease slows dramatically
- Elbow indicates optimal balance between fit and complexity

Limitation: Clear elbows rare with real consumer data!

Elbow at k=3

Figure 8: Idealized scree plot showing clear elbow at $k=3$ segments. In practice, consumer data rarely produces such distinct elbows.

6.2 Stability Analysis: The Gold Standard

Why Stability Matters

Consumer datasets typically lack natural, well-separated clusters. Multiple segmentation solutions may be statistically equivalent.

Stability analysis answers:

- Are extracted segments **reproducible**?
- Which k produces most **stable** solutions?
- Are individual segments **robust** to sampling variation?

6.2.1 Global Stability Analysis

Bootstrap Stability Procedure

- 1: Generate b pairs of bootstrap samples (2b total)
- 2: **for** each $k = 2, 3, \dots, k_{\max}$ **do**
- 3: **for** each pair of bootstrap samples **do**
- 4: Cluster both samples into k segments
- 5: Compute adjusted Rand index between the two solutions
- 6: **end for**

```

7:     Store  $b$  adjusted Rand indices for this  $k$ 
8: end for
9: Create boxplots of adjusted Rand indices vs.  $k$ 
10: Interpretation:
11:     High values ( $\approx 1$ ): Reproducible structure
12:     Low values ( $\approx 0$ ): Artificial construction
13: return Recommended  $k$  with highest median stability

```

Adjusted Rand Index

Measures similarity between two partitions, correcting for chance:

$$\text{ARI} = \frac{\text{Index} - \text{Expected Index}}{\text{Max Index} - \text{Expected Index}} \quad (15)$$

Properties:

- ARI = 1: Perfect agreement
- ARI = 0: Agreement by chance
- ARI < 0 : Less agreement than expected by chance

6.2.2 Python Implementation: Stability Analysis

```

1 import numpy as np
2 from sklearn.cluster import KMeans
3 from sklearn.metrics import adjusted_rand_score
4 import matplotlib.pyplot as plt
5
6 def bootstrap_stability_analysis(data, k_range, n_bootstrap_pairs=50):
7     """
8         Perform global stability analysis using bootstrap resampling.
9     """
10    n = len(data)
11    ari_results = {k: [] for k in k_range}
12
13    print("Running bootstrap stability analysis...")
14    for pair_idx in range(n_bootstrap_pairs):
15        if pair_idx % 10 == 0:
16            print(f"  Bootstrap pair {pair_idx+1}/{n_bootstrap_pairs}")
17
18        # Generate pair of bootstrap samples
19        indices1 = np.random.choice(n, size=n, replace=True)
20        indices2 = np.random.choice(n, size=n, replace=True)
21
22        sample1 = data[indices1]
23        sample2 = data[indices2]
24
25        for k in k_range:
26            # Cluster both samples
27            kmeans1 = KMeans(n_clusters=k, n_init=10, random_state=42)

```

```

28         kmeans2 = KMeans(n_clusters=k, n_init=10, random_state=42)
29
30     labels1 = kmeans1.fit_predict(sample1)
31     labels2 = kmeans2.fit_predict(sample2)
32
33     # Compute adjusted Rand index
34     ari = adjusted_rand_score(labels1, labels2)
35     ari_results[k].append(ari)
36
37 return ari_results
38
39 # Example usage
40 np.random.seed(42)
41
42 # Generate data with 3 natural segments
43 seg1 = np.random.normal([2, 2], [0.5, 0.5], (200, 2))
44 seg2 = np.random.normal([5, 5], [0.5, 0.5], (200, 2))
45 seg3 = np.random.normal([8, 2], [0.5, 0.5], (200, 2))
46 data = np.vstack([seg1, seg2, seg3])
47
48 # Run stability analysis
49 k_range = range(2, 9)
50 ari_results = bootstrap_stability_analysis(data, k_range, n_bootstrap_pairs=50)
51
52 # Create stability boxplot
53 fig, ax = plt.subplots(figsize=(10, 6))
54
55 positions = list(k_range)
56 ari_data = [ari_results[k] for k in k_range]
57
58 bp = ax.boxplot(ari_data, positions=positions, widths=0.6, patch_artist=True,
59                  boxprops=dict(facecolor='lightblue', edgecolor='blue'),
60                  medianprops=dict(color='red', linewidth=2),
61                  whiskerprops=dict(color='blue'),
62                  capprops=dict(color='blue'))
63
64 ax.set_xlabel('Number of Segments (k)', fontsize=12)
65 ax.set_ylabel('Adjusted Rand Index', fontsize=12)
66 ax.set_title('Global Stability Analysis', fontsize=14, fontweight='bold')
67 ax.set_ylim([0, 1.05])
68 ax.grid(True, alpha=0.3)
69 ax.axhline(y=0.8, color='green', linestyle='--', alpha=0.5,
70             label='High Stability Threshold')
71
72 # Highlight optimal k
73 medians = [np.median(ari_results[k]) for k in k_range]
74 optimal_k = list(k_range)[np.argmax(medians)]
75 ax.axvline(x=optimal_k, color='red', linestyle='--', linewidth=2,
76             label=f'Optimal k={optimal_k}')
77
78 ax.legend()
79 plt.tight_layout()
80 plt.show()
81
82 print(f"\nOptimal number of segments: {optimal_k}")
83 print("\nMedian ARI by k:")

```

```
84 for k in k_range:  
85     print(f"  k={k}: {np.median(ari_results[k]):.3f}")
```

7 Practical Recommendations

Algorithm Selection Guidelines

1. Start with exploration:

- Hierarchical clustering with different linkages
- k-means with range of k values
- Self-organizing maps for visualization

2. Assess data structure:

- Run global stability analysis
- Generate scree plots (with caution)
- Inspect dendograms

3. Refine with advanced methods:

- Finite mixture models for probabilistic segments
- Bagged clustering for niche segment detection
- VSBD for high-dimensional binary data

4. Validate thoroughly:

- Segment-level stability analysis
- Business interpretation and actionability
- Cross-validation with holdout samples

Critical Questions for Practitioners

Before finalizing a segmentation solution, ask:

1. **Stability:** Do segments replicate across bootstrap samples?
2. **Interpretability:** Can we clearly describe each segment?
3. **Actionability:** Can we target these segments differently?
4. **Size viability:** Are segments large enough to be profitable?
5. **Accessibility:** Can we reach these segments through marketing channels?

8 Conclusion

Key Takeaways

Essential Insights:

1. **No universally best algorithm:** Algorithm choice shapes results
2. **Consumer data lacks natural structure:** Expect reproducible or constructive segmentation, not natural clusters
3. **Stability is paramount:** Reproducibility across samples indicates meaningful segments
4. **Domain knowledge essential:** Statistical methods guide, but business insight determines final selection
5. **Multiple solutions valid:** Explore alternatives, document rationale for final choice

References

- [1] Dolnicar, S., Grün, B., and Leisch, F. (2018). *Market Segmentation Analysis: Understanding It, Doing It, and Making It Useful*. Springer.
- [2] Hennig, C. (2015). What are the true clusters? *Pattern Recognition Letters*, 64, 53-62.
- [3] Wedel, M., and Kamakura, W.A. (2000). *Market Segmentation: Conceptual and Methodological Foundations*. Kluwer Academic Publishers.
- [4] Fraley, C., and Raftery, A.E. (2002). Model-based clustering, discriminant analysis, and density estimation. *Journal of the American Statistical Association*, 97(458), 611-631.
- [5] Kohonen, T. (2001). *Self-Organizing Maps*. 3rd ed., Springer.

9 Case Studies: From Theory to Practice

9.1 Case Study 1: Tourist Vacation Activities

Business Context

A tourism organization wants to segment tourists based on their vacation activity patterns to design targeted marketing campaigns and service offerings.

Dataset: 7 tourists, 3 vacation activities (Beach, Action, Culture)

Objective: Identify natural groupings of tourists with similar activity preferences

Approach: Distance-based hierarchical clustering with multiple linkage methods

9.1.1 Data Structure and Initial Exploration

Table 4: Tourist Vacation Activities Dataset

Tourist	Beach (%)	Action (%)	Culture (%)
Anna	100	0	0
Bill	100	0	0
Frank	60	40	0
Julia	70	0	30
Maria	80	0	20
Michael	0	90	10
Tom	50	20	30

Initial Observations

Examining the raw data reveals:

- **Anna & Bill:** Identical profiles (100% beach orientation)
- **Michael:** Distinct outlier (no beach interest, high action)
- **Julia & Maria:** Similar beach+culture mix
- **Frank & Tom:** Mixed activity patterns

Expected outcome: 2-3 meaningful segments depending on granularity desired

9.1.2 Distance Matrix Analysis

```

1 import numpy as np
2 import pandas as pd
3 from scipy.spatial.distance import pdist, squareform
4 from scipy.cluster.hierarchy import dendrogram, linkage
5 import matplotlib.pyplot as plt
6
7 # Create dataset
8 data = pd.DataFrame({
9     'Beach': [100, 100, 60, 70, 80, 0, 50],
10    'Action': [0, 0, 40, 0, 0, 90, 20],
11    'Culture': [0, 0, 0, 30, 20, 10, 30]
12 }, index=['Anna', 'Bill', 'Frank', 'Julia', 'Maria', 'Michael', 'Tom'])
13
14 # Calculate Manhattan distance matrix
15 dist_manhattan = pdist(data, metric='cityblock')
16 dist_matrix = squareform(dist_manhattan)
17
18 # Create distance matrix DataFrame
19 dist_df = pd.DataFrame(dist_matrix,
20                         index=data.index,
21                         columns=data.index)
22
23 print("Manhattan Distance Matrix:")

```

```

24 print(dist_df.astype(int))
25
26 # Output:
27 #      Anna  Bill  Frank  Julia  Maria  Michael  Tom
28 # Anna    0     0     80     60     40     200    100
29 # Bill    0     0     80     60     40     200    100
30 # Frank   80    80     0     80     80     120    60
31 # Julia   60    60     80     0     20     180    40
32 # Maria   40    40     80     20     0     180    60
33 # Michael 200   200    120    180    180     0    140
34 # Tom    100   100     60     40     60    140     0

```

Distance Matrix Insights

Key Patterns:

- Anna-Bill distance = 0 (identical twins)
- Julia-Maria distance = 20 (close neighbors)
- Michael has large distances to everyone (outlier status)
- Tom shows moderate distances to most (bridge consumer)

9.1.3 Hierarchical Clustering Analysis

```

1 # Perform hierarchical clustering with different linkages
2 linkages = ['single', 'complete', 'average', 'ward']
3 results = {}
4
5 fig, axes = plt.subplots(2, 2, figsize=(14, 10))
6 axes = axes.ravel()
7
8 for idx, method in enumerate(linkages):
9     # Compute linkage
10    Z = linkage(dist_manhattan, method=method)
11    results[method] = Z
12
13    # Plot dendrogram
14    ax = axes[idx]
15    dendrogram(Z, labels=data.index, ax=ax)
16    ax.set_title(f'{method.capitalize()} Linkage', fontsize=14,
17                 fontweight='bold')
17    ax.set_xlabel('Tourist')
18    ax.set_ylabel('Distance')
19
20    # Add cut line for k=3
21    if method in ['complete', 'average']:
22        ax.axhline(y=100, color='red', linestyle='--',
23                    label='k=3 cut', linewidth=2)
24        ax.legend()
25
26 plt.tight_layout()
27 plt.show()

```

```

28
29 # Extract segments for complete linkage (k=3)
30 from scipy.cluster.hierarchy import fcluster
31
32 segments_complete = fcluster(results['complete'], t=3, criterion='maxclust')
33 data['Segment_Complete'] = segments_complete
34
35 print("\nSegment Assignment (Complete Linkage, k=3):")
36 for segment in range(1, 4):
37     members = data[data['Segment_Complete'] == segment].index.tolist()
38     print(f"Segment {segment}: {', '.join(members)}")

```

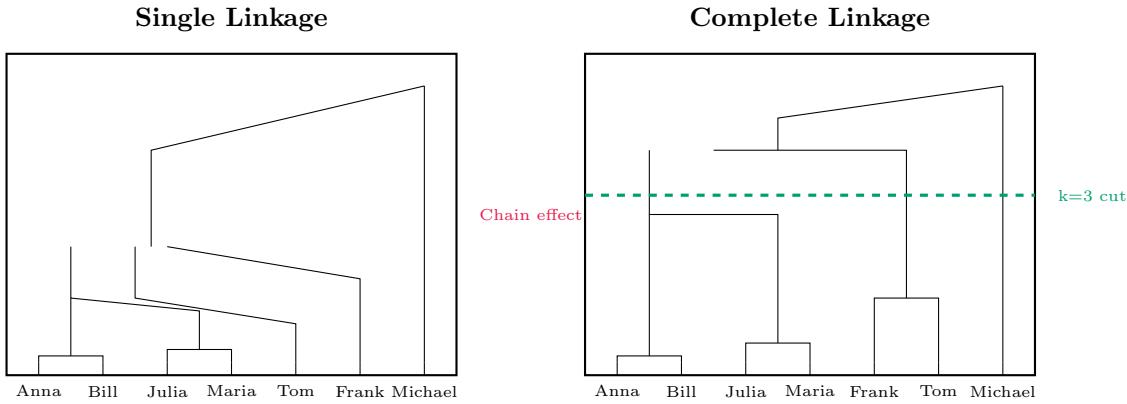


Figure 9: Comparison of single and complete linkage hierarchical clustering on tourist data. Single linkage shows chaining effect, while complete linkage produces more compact, interpretable segments.

9.1.4 Segment Profiling

```

1 # Calculate segment profiles
2 segment_profiles = data.groupby('Segment_Complete')[['Beach', 'Action',
   ↵ 'Culture']].mean()
3
4 print("\nSegment Profiles (Complete Linkage, k=3):")
5 print(segment_profiles.round(1))
6
7 # Visualize segment profiles
8 fig, ax = plt.subplots(figsize=(10, 6))
9
10 x = np.arange(len(segment_profiles.columns))
11 width = 0.25
12
13 for i, segment in enumerate([1, 2, 3]):
14     segment_data = segment_profiles.loc[segment]
15     ax.bar(x + i*width, segment_data, width,
16             label=f'Segment {segment}', alpha=0.8)
17
18 # Add population mean
19 pop_mean = data[['Beach', 'Action', 'Culture']].mean()
20 ax.plot(x + width, pop_mean, 'ro-', markersize=10, linewidth=2,

```

```

21     label='Population Mean')
22
23 ax.set_xlabel('Activity', fontsize=12)
24 ax.set_ylabel('Percentage of Time', fontsize=12)
25 ax.set_title('Segment Activity Profiles', fontsize=14, fontweight='bold')
26 ax.set_xticks(x + width)
27 ax.set_xticks(segment_profiles.columns)
28 ax.legend()
29 ax.grid(axis='y', alpha=0.3)
30
31 plt.tight_layout()
32 plt.show()

```

Segment Interpretation Naming

Segment 1: "Beach Lovers" (Anna, Bill, Julia, Maria)

- High beach orientation (85% average)
- Minimal action activities
- Some cultural interest (Julia, Maria)
- **Marketing strategy:** Beach resort packages, relaxation focus

Segment 2: "Mixed Activity Seekers" (Frank, Tom)

- Balanced activity portfolio
- Moderate beach, action, and culture
- **Marketing strategy:** Diverse destination offerings, flexibility

Segment 3: "Action Enthusiasts" (Michael)

- No beach interest
- Dominant action focus (90%)
- Niche segment (single member in sample)
- **Marketing strategy:** Adventure tourism, extreme sports

9.2 Case Study 2: Austrian Winter Vacation Activities

Business Context

An Austrian tourism board seeks to segment winter tourists based on 27 different vacation activities to develop targeted marketing campaigns for different tourist types.

Dataset: 2,961 tourists, 27 binary activities

Challenge: High-dimensional binary data with potential niche segments

Approach: Bagged clustering to identify both major and niche segments

9.2.1 Dataset Overview

The dataset includes diverse winter activities: - **Winter sports:** Alpine skiing, cross-country skiing, snowboarding, ice-skating - **Wellness:** Spa visits, health facilities, pool/sauna - **Culture:** Museums, theater/opera, concerts - **Recreation:** Hiking, walks, excursions, relaxation - **Entertainment:** Evening activities, discos/bars, shopping

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.cluster import KMeans
4 from sklearn.preprocessing import StandardScaler
5 from scipy.cluster.hierarchy import dendrogram, linkage, fcluster
6 import matplotlib.pyplot as plt
7
8 # Simulate winter activities data (binary)
9 np.random.seed(42)
10 n_tourists = 2961
11 n_activities = 27
12
13 # Create realistic activity patterns
14 # Segment 1: Alpine skiers (35%)
15 n_seg1 = int(0.35 * n_tourists)
16 seg1 = np.random.binomial(1, 0.8, (n_seg1, 5)) # High winter sports
17 seg1 = np.hstack([seg1, np.random.binomial(1, 0.3, (n_seg1, 22))])
18
19 # Segment 2: Health tourists (11%)
20 n_seg2 = int(0.11 * n_tourists)
21 seg2 = np.random.binomial(1, 0.2, (n_seg2, 27))
22 seg2[:, [9, 10, 26]] = np.random.binomial(1, 0.7, (n_seg2, 3)) # Spa, health,
23   ↪ pool
24
25 # Segment 3: Cultural tourists (16%)
26 n_seg3 = int(0.16 * n_tourists)
27 seg3 = np.random.binomial(1, 0.3, (n_seg3, 27))
28 seg3[:, [19, 20, 21, 22]] = np.random.binomial(1, 0.6, (n_seg3, 4)) # Culture
29
30 # Segment 4: Mixed activities (38%)
31 n_seg4 = n_tourists - n_seg1 - n_seg2 - n_seg3
32 seg4 = np.random.binomial(1, 0.4, (n_seg4, 27))
33
34 # Combine
35 data = np.vstack([seg1, seg2, seg3, seg4])
36 true_labels = np.array([0]*n_seg1 + [1]*n_seg2 + [2]*n_seg3 + [3]*n_seg4)
37
38 activity_names = [
39     'alpine_skiing', 'cross_country', 'snowboarding', 'ice_skating',
40     ↪ 'ski_touring',
41     'sleigh_riding', 'tennis', 'horseback_riding', 'going_to_spa',
42     ↪ 'health_facilities',
43     'hiking', 'walks', 'organized_excursions', 'excursions', 'relaxing',
44     'evening_out', 'discos_bars', 'shopping', 'sight_seeing', 'museums',
45     'theater_opera', 'heurigen', 'concerts', 'tyrolean_evenings',
46     ↪ 'local_events',
47     'pool_sauna', 'other'
48

```

```

44 ]
45
46 df = pd.DataFrame(data, columns=activity_names[:27])

```

9.2.2 Bagged Clustering Implementation

```

1 def bagged_clustering(data, base_k=10, n_bootstrap=50):
2     """
3         Perform bagged clustering on binary activity data.
4     """
5     n_samples = len(data)
6     all_centroids = []
7
8     print(f"Performing bagged clustering with {n_bootstrap} bootstrap
9         samples...")
10
11    for b in range(n_bootstrap):
12        # Bootstrap sample
13        indices = np.random.choice(n_samples, size=n_samples, replace=True)
14        bootstrap_data = data[indices]
15
16        # k-means with base_k clusters
17        kmeans = KMeans(n_clusters=base_k, n_init=10, random_state=b)
18        kmeans.fit(bootstrap_data)
19
20        # Store centroids
21        all_centroids.append(kmeans.cluster_centers_)
22
23        if (b + 1) % 10 == 0:
24            print(f"  Completed {(b+1)}/{n_bootstrap} bootstrap samples")
25
26    # Stack all centroids
27    all_centroids = np.vstack(all_centroids)
28    print(f"\nCollected {len(all_centroids)} centroids")
29
30    # Hierarchical clustering on centroids
31    linkage_matrix = linkage(all_centroids, method='average')
32
33    return linkage_matrix, all_centroids
34
35 # Run bagged clustering
36 linkage_matrix, all_centroids = bagged_clustering(df.values,
37                                                 base_k=10,
38                                                 n_bootstrap=50)
39
40 # Plot dendrogram
41 plt.figure(figsize=(12, 6))
42 dendrogram(linkage_matrix, no_labels=True)
43 plt.xlabel('Centroid Index')
44 plt.ylabel('Distance')
45 plt.title('Bagged Clustering Dendrogram\nAustrian Winter Activities',
46             fontsize=14, fontweight='bold')
47 plt.axhline(y=2.5, color='red', linestyle='--', linewidth=2, label='k=5 cut')
48 plt.legend()

```

```

48 plt.tight_layout()
49 plt.show()
50
51 # Extract k=5 segments
52 final_k = 5
53 centroid_labels = fcluster(linkage_matrix, t=final_k, criterion='maxclust')
54
55 # Get final segment centers
56 final_centers = np.array([all_centroids[centroid_labels == j].mean(axis=0)
57                           for j in range(1, final_k+1)])
58
59 # Assign original tourists to segments
60 from scipy.spatial.distance import cdist
61 distances = cdist(df.values, final_centers)
62 tourist_segments = np.argmin(distances, axis=1)
63
64 df['Segment'] = tourist_segments

```

9.2.3 Segment Profiles and Business Insights

```

1 # Calculate segment activity probabilities
2 segment_profiles = df.groupby('Segment').mean() * 100 # Convert to percentages
3
4 # Visualize top activities per segment
5 fig, axes = plt.subplots(1, 5, figsize=(18, 6))
6
7 for seg in range(5):
8     ax = axes[seg]
9     profile = segment_profiles.iloc[seg].sort_values(ascending=False)[:10]
10
11    ax.bars(range(len(profile)), profile.values, color=f'C{seg}')
12    ax.set_yticks(range(len(profile)))
13    ax.set_yticklabels(profile.index, fontsize=8)
14    ax.set_xlabel('Activity Rate (%)')
15    ax.set_title(f'Segment {seg+1}\n{(df["Segment"]==seg).sum()} tourists')
16    ax.invert_yaxis()
17    ax.grid(axis='x', alpha=0.3)
18
19 plt.tight_layout()
20 plt.show()
21
22 # Segment sizes
23 print("\nSegment Sizes:")
24 print(df['Segment'].value_counts().sort_index())
25 print("\nSegment Size Distribution:")
26 print(df['Segment'].value_counts(normalize=True).sort_index() * 100)

```

10 Advanced Stability Analysis

10.1 Segment-Level Stability Within Solutions (SLSw)

Traditional global stability analysis assesses entire segmentation solutions. However, organizations typically target only one or a few segments, not the complete solution. Segment-

level stability within solutions (SLSw) addresses this limitation.

The SLSw Principle

Segment-level stability within solutions (SLSw) answers:

- Which **individual segments** are most reproducible?
- Can we identify a single highly stable segment for targeting?
- Does the solution contain attractive niche segments?

Key insight: A globally unstable solution may contain one highly stable, valuable segment worth targeting.

10.1.1 SLSw Algorithm

Segment-Level Stability Within Solutions

Based on Hennig (2007) and Dolnicar & Leisch (2017)

```

1: Extract partition with  $k$  segments:  $S_1, \dots, S_k$ 
2: Draw  $b$  bootstrap samples (typical:  $b = 100$ )
3: for each bootstrap sample  $i = 1$  to  $b$  do
4:   Extract  $k$  segments:  $S_1^{(i)}, \dots, S_k^{(i)}$ 
5: end for
6: for each original segment  $h = 1$  to  $k$  do
7:   for each bootstrap sample  $i = 1$  to  $b$  do
8:     Find bootstrap segment  $S_j^{(i)}$  with maximum Jaccard similarity to  $S_h$ 
9:     Store Jaccard index:  $J_h^{(i)}$ 
10:    end for
11:    Compute distribution of  $\{J_h^{(1)}, \dots, J_h^{(b)}\}$ 
12:  end for
13: return Boxplots of Jaccard indices for each segment

```

Jaccard Index for Segment Similarity

For two segments A and B :

$$J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|} \quad (16)$$

where $|A \cap B|$ is number of consumers in both segments.

Interpretation:

- $J = 1$: Perfect agreement (identical segments)
- $J = 0.8$: High stability (80% overlap)
- $J < 0.5$: Low stability (artificial segment)

10.1.2 Python Implementation: SLSw Analysis

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.cluster import KMeans
4 import matplotlib.pyplot as plt
5
6 def jaccard_index(set1, set2):
7     """Calculate Jaccard index between two sets."""
8     intersection = len(set(set1) & set(set2))
9     union = len(set(set1) | set(set2))
10    return intersection / union if union > 0 else 0
11
12 def segment_level_stability_within(data, k, n_bootstrap=100, random_state=42):
13     """
14     Compute segment-level stability within solutions (SLSw).
15
16     Returns: Dictionary with Jaccard indices for each segment
17     """
18     np.random.seed(random_state)
19     n_samples = len(data)
20
21     # Fit original solution
22     kmeans_orig = KMeans(n_clusters=k, n_init=50, random_state=random_state)
23     labels_orig = kmeans_orig.fit_predict(data)
24
25     # Store Jaccard indices for each segment
26     jaccard_results = {seg: [] for seg in range(k)}
27
28     print(f"Computing SLSw for k={k} with {n_bootstrap} bootstrap samples...")
29
30     for b in range(n_bootstrap):
31         # Bootstrap sample
32         indices = np.random.choice(n_samples, size=n_samples, replace=True)
33         bootstrap_data = data[indices]
34
35         # Cluster bootstrap sample
36         kmeans_boot = KMeans(n_clusters=k, n_init=10, random_state=b)
37         labels_boot = kmeans_boot.fit_predict(bootstrap_data)
38
39         # For each original segment, find best matching bootstrap segment
40         for orig_seg in range(k):
41             # Original segment members (in bootstrap sample indices)
42             orig_members = np.where(labels_orig == orig_seg)[0]
43             orig_in_boot = np.array([np.where(indices == m)[0][0]
44                                    if m in indices else -1
45                                    for m in orig_members])
46             orig_in_boot = orig_in_boot[orig_in_boot >= 0]
47
48             # Find best matching bootstrap segment
49             max_jaccard = 0
50             for boot_seg in range(k):
51                 boot_members = np.where(labels_boot == boot_seg)[0]
52                 jaccard = jaccard_index(orig_in_boot, boot_members)
53                 max_jaccard = max(max_jaccard, jaccard)

```

```

54         jaccard_results[orig_seg].append(max_jaccard)
55
56     if (b + 1) % 20 == 0:
57         print(f"  Completed {(b+1)}/{n_bootstrap} bootstrap samples")
58
59     return jaccard_results
60
61
62 # Example usage
63 np.random.seed(42)
64
65 # Generate data with 3 segments of different stability
66 stable_seg = np.random.normal([2, 2], [0.3, 0.3], (200, 2))
67 moderate_seg = np.random.normal([5, 5], [0.8, 0.8], (200, 2))
68 unstable_seg = np.random.normal([8, 2], [1.5, 1.5], (200, 2))
69
70 data = np.vstack([stable_seg, moderate_seg, unstable_seg])
71
72 # Compute SLSw
73 slsw_results = segment_level_stability_within(data, k=3, n_bootstrap=100)
74
75 # Visualize
76 fig, ax = plt.subplots(figsize=(10, 6))
77
78 positions = [1, 2, 3]
79 slsw_data = [slsw_results[i] for i in range(3)]
80
81 bp = ax.boxplot(slsw_data, positions=positions, widths=0.6, patch_artist=True,
82                  boxprops=dict(facecolor='lightblue', edgecolor='blue'),
83                  medianprops=dict(color='red', linewidth=2),
84                  whiskerprops=dict(color='blue'),
85                  capprops=dict(color='blue'))
86
87 ax.set_xlabel('Segment Number', fontsize=12)
88 ax.set_ylabel('Jaccard Coefficient', fontsize=12)
89 ax.set_title('Segment-Level Stability Within Solutions (SLSw)',
90               fontsize=14, fontweight='bold')
91 ax.set_ylim([0, 1.05])
92 ax.set_xticks(positions)
93 ax.grid(True, alpha=0.3, axis='y')
94 ax.axhline(y=0.8, color='green', linestyle='--', alpha=0.5,
95             label='High Stability Threshold')
96
97 # Add median values as text
98 for i, pos in enumerate(positions):
99     median_val = np.median(slsw_results[i])
100    ax.text(pos, median_val + 0.05, f'{median_val:.3f}',
101            ha='center', fontweight='bold')
102
103 ax.legend()
104 plt.tight_layout()
105 plt.show()
106
107 # Print statistics
108 print("\nSLSw Statistics:")
109 for seg in range(3):

```

```

110     print(f"\nSegment {seg+1}:")
111     print(f"  Median Jaccard: {np.median(slsw_results[seg]):.3f}")
112     print(f"  Mean Jaccard: {np.mean(slsw_results[seg]):.3f}")
113     print(f"  Std Dev: {np.std(slsw_results[seg]):.3f}")

```

10.2 Segment-Level Stability Across Solutions (SLSA)

The SLSA Principle

Segment-level stability across solutions (SLSA) tracks how segments evolve as k increases:

Questions addressed:

- Do certain segments persist across different k values?
- Which segments are "natural" vs. artificially created?
- How are segments split or merged as k changes?

Interpretation:

- **Thick horizontal lines:** Stable segments that persist
- **Branching patterns:** Segments splitting with increasing k
- **Multiple inputs:** Artificial segments created by merging

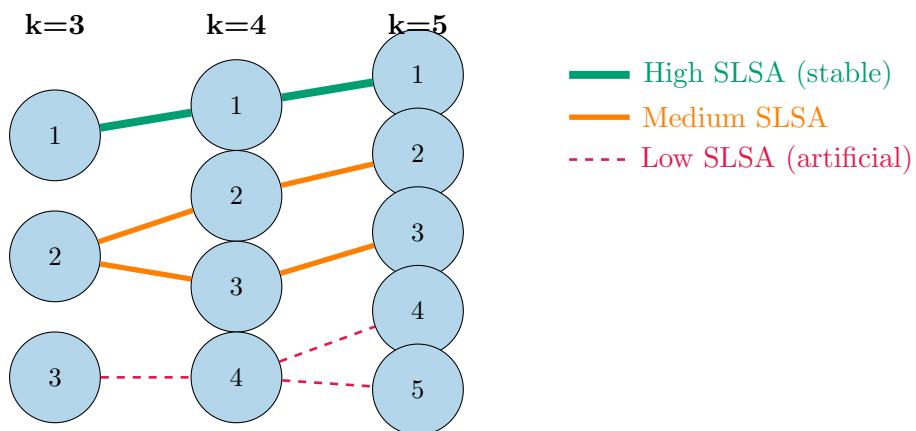


Figure 10: Segment-level stability across solutions (SLSA) visualization. Segment 1 remains highly stable across k values (thick green lines). Segment 2 progressively splits (orange). Segment 3 is unstable and artificial (red dashed).

10.2.1 SLSA with Entropy-Based Stability

Entropy-Based SLSA Measure

For segment S_i in solution with k_i segments, recruiting from segments in solution with k_{i-1} segments:

Let p_j = proportion of consumers in S_i from segment j in previous solution

Entropy:

$$H(S_i) = - \sum_{j=1}^{k_{i-1}} p_j \log(p_j) \quad (17)$$

Normalized SLSA stability:

$$\text{SLSA}(S_i) = 1 - \frac{H(S_i)}{\log(k_{i-1})} \quad (18)$$

Interpretation:

- SLSA = 1: All members from single parent segment (maximum stability)
- SLSA = 0: Members equally distributed from all parent segments (artificial)
- SLSA > 0.8: Highly stable segment

10.3 Response Style Detection

Response Bias in Survey Data

Survey-based segmentation faces a critical challenge: **response style bias**

Common response styles:

- **Acquiescence:** Tendency to agree with all statements
- **Disacquiescence:** Tendency to disagree with all statements
- **Midpoint:** Always selecting neutral/middle options
- **Extreme response:** Only using endpoint scales

Consequence: Segments may reflect response patterns, not true preferences!

10.3.1 Detecting Response Style Segments

Response Style Indicators

A segment likely reflects response style if:

1. **High agreement across all variables**
 - Acquiescence: >80% agreement on most items
 - Uniform pattern regardless of content
2. **High stability but questionable profile**

- Segment appears stable in SLSw/SLSA analysis
- Profile lacks coherent narrative
- Members agree with contradictory statements

3. Correlation with demographic "yes-sayers"

- Older respondents often show acquiescence
- Lower education sometimes correlates with midpoint bias

Action: Exclude response style segments from targeting consideration

10.3.2 Python Implementation: Response Style Detection

```

1 import numpy as np
2 import pandas as pd
3 from sklearn.cluster import KMeans
4 import matplotlib.pyplot as plt
5
6 def detect_response_style(segment_data, threshold_agreement=0.8):
7     """
8         Detect potential response style bias in a segment.
9
10    Parameters:
11        - segment_data: DataFrame with binary/Likert responses
12        - threshold_agreement: Minimum proportion for acquiescence flag
13
14    Returns: Dictionary with response style indicators
15    """
16    n_variables = segment_data.shape[1]
17
18    # Calculate mean agreement per variable
19    agreement_rates = segment_data.mean()
20
21    # Check for acquiescence (high agreement across board)
22    high_agreement_count = (agreement_rates > threshold_agreement).sum()
23    acquiescence_ratio = high_agreement_count / n_variables
24
25    # Check for variance (low variance suggests response style)
26    variances = segment_data.var()
27    low_variance_count = (variances < 0.1).sum()
28    low_variance_ratio = low_variance_count / n_variables
29
30    # Overall response style score
31    response_style_score = (acquiescence_ratio + low_variance_ratio) / 2
32
33    results = {
34        'acquiescence_ratio': acquiescence_ratio,
35        'low_variance_ratio': low_variance_ratio,
36        'response_style_score': response_style_score,
37        'is_response_style': response_style_score > 0.6,
38        'mean_agreement': agreement_rates.mean(),
39        'std_agreement': agreement_rates.std()

```

```

40     }
41
42     return results
43
44 # Example: Simulate survey data
45 np.random.seed(42)
46
47 # Segment 1: True preference pattern (varied responses)
48 true_segment = np.random.binomial(1, [0.2, 0.8, 0.3, 0.7, 0.4], (100, 5))
49
50 # Segment 2: Acquiescence bias (agrees with everything)
51 acquiescence_segment = np.random.binomial(1, 0.85, (100, 5))
52
53 # Segment 3: Midpoint bias (moderate on everything)
54 midpoint_segment = np.random.binomial(1, 0.5, (100, 5))
55
56 segments_data = [
57     pd.DataFrame(true_segment, columns=[f'Q{i+1}' for i in range(5)]),
58     pd.DataFrame(acquiescence_segment, columns=[f'Q{i+1}' for i in range(5)]),
59     pd.DataFrame(midpoint_segment, columns=[f'Q{i+1}' for i in range(5)])
60 ]
61
62 # Analyze each segment
63 print("Response Style Analysis:")
64 print("=" * 60)
65
66 for i, seg_data in enumerate(segments_data, 1):
67     print(f"\nSegment {i}:")
68     results = detect_response_style(seg_data)
69
70     print(f"    Acquiescence Ratio: {results['acquiescence_ratio']:.3f}")
71     print(f"    Low Variance Ratio: {results['low_variance_ratio']:.3f}")
72     print(f"    Response Style Score: {results['response_style_score']:.3f}")
73     print(f"    Response Style Detected: {results['is_response_style']} ")
74     print(f"    Mean Agreement: {results['mean_agreement']:.3f}")
75     print(f"    Std Agreement: {results['std_agreement']:.3f}")
76
77 # Visualize patterns
78 fig, axes = plt.subplots(1, 3, figsize=(15, 4))
79
80 segment_names = ['True Preference', 'Acquiescence Bias', 'Midpoint Bias']
81
82 for i, (seg_data, ax) in enumerate(zip(segments_data, axes)):
83     agreement_rates = seg_data.mean()
84
85     ax.bar(range(5), agreement_rates, color=f'C{i}', alpha=0.7)
86     ax.set_xlabel('Question')
87     ax.set_ylabel('Agreement Rate')
88     ax.set_title(segment_names[i])
89     ax.set_ylim([0, 1])
90     ax.axhline(y=0.8, color='red', linestyle='--', alpha=0.5, label='Threshold')
91     ax.set_xticks(range(5))
92     ax.set_xticklabels([f'Q{j+1}' for j in range(5)])
93     ax.legend()
94     ax.grid(axis='y', alpha=0.3)
95

```

```
96 plt.tight_layout()
97 plt.show()
```

10.4 Handling Mixed Data Types

Real consumer datasets often contain mixed variable types: binary activities, ordinal ratings, metric measurements.

Mixed Data Challenges

Common variable types in consumer data:

- **Binary:** Purchased/not purchased, activity participation
- **Ordinal:** Likert scales, satisfaction ratings
- **Metric:** Spending amounts, age, income
- **Nominal:** Categorical preferences (brand, color)

Problem: Standard distance measures assume all variables are same type

10.4.1 Gower Distance for Mixed Data

Gower Distance Formula

For consumers i and j with p mixed-type variables:

$$d_{Gower}(i, j) = \frac{\sum_{k=1}^p \delta_{ijk} \cdot d_{ijk}}{\sum_{k=1}^p \delta_{ijk}} \quad (19)$$

where:

- $\delta_{ijk} = 0$ if variable k is missing for i or j , else 1
- d_{ijk} depends on variable type:

For binary/nominal: $d_{ijk} = \begin{cases} 0 & \text{if } x_{ik} = x_{jk} \\ 1 & \text{otherwise} \end{cases}$

For ordinal/metric: $d_{ijk} = \frac{|x_{ik} - x_{jk}|}{\text{range}_k}$

```
1 from sklearn.preprocessing import LabelEncoder
2 from scipy.spatial.distance import pdist, squareform
3 import numpy as np
4
5 def gower_distance(X, categorical_features=None):
6     """
7         Compute Gower distance for mixed data types.
8
9     Parameters:
10        - X: DataFrame with mixed types
11        - categorical_features: List of categorical column names
```

```

12
13     Returns: Distance matrix
14     """
15
16     X = X.copy()
17     n = len(X)
18
19     if categorical_features is None:
20         categorical_features = X.select_dtypes(include=['object',
21             'category']).columns
22
23     numerical_features = [c for c in X.columns if c not in categorical_features]
24
25
26     for i in range(n):
27         for j in range(i+1, n):
28             distance = 0
29             valid_features = 0
30
31             # Categorical features
32             for feat in categorical_features:
33                 if pd.notna(X.iloc[i][feat]) and pd.notna(X.iloc[j][feat]):
34                     distance += 0 if X.iloc[i][feat] == X.iloc[j][feat] else 1
35                     valid_features += 1
36
37             # Numerical features (normalized)
38             for feat in numerical_features:
39                 if pd.notna(X.iloc[i][feat]) and pd.notna(X.iloc[j][feat]):
40                     feat_range = X[feat].max() - X[feat].min()
41                     if feat_range > 0:
42                         distance += abs(X.iloc[i][feat] - X.iloc[j][feat]) /
43                             feat_range
44                         valid_features += 1
45
46             distance_matrix[i, j] = distance / valid_features if valid_features
47             > 0 else 0
48             distance_matrix[j, i] = distance_matrix[i, j]
49
50     return distance_matrix
51
52 # Example usage with mixed data
53 data_mixed = pd.DataFrame({
54     'age': [25, 30, 28, 45, 50],
55     'income': [30000, 45000, 35000, 70000, 65000],
56     'purchased': [1, 0, 1, 0, 1],
57     'preferred_brand': ['A', 'B', 'A', 'C', 'C']
58 })
59
60 dist_matrix = gower_distance(data_mixed,
61     categorical_features=['preferred_brand'])
62 print("Gower Distance Matrix:")
63 print(dist_matrix)

```


11 Complete Market Segmentation Workflow

11.1 End-to-End Process Overview

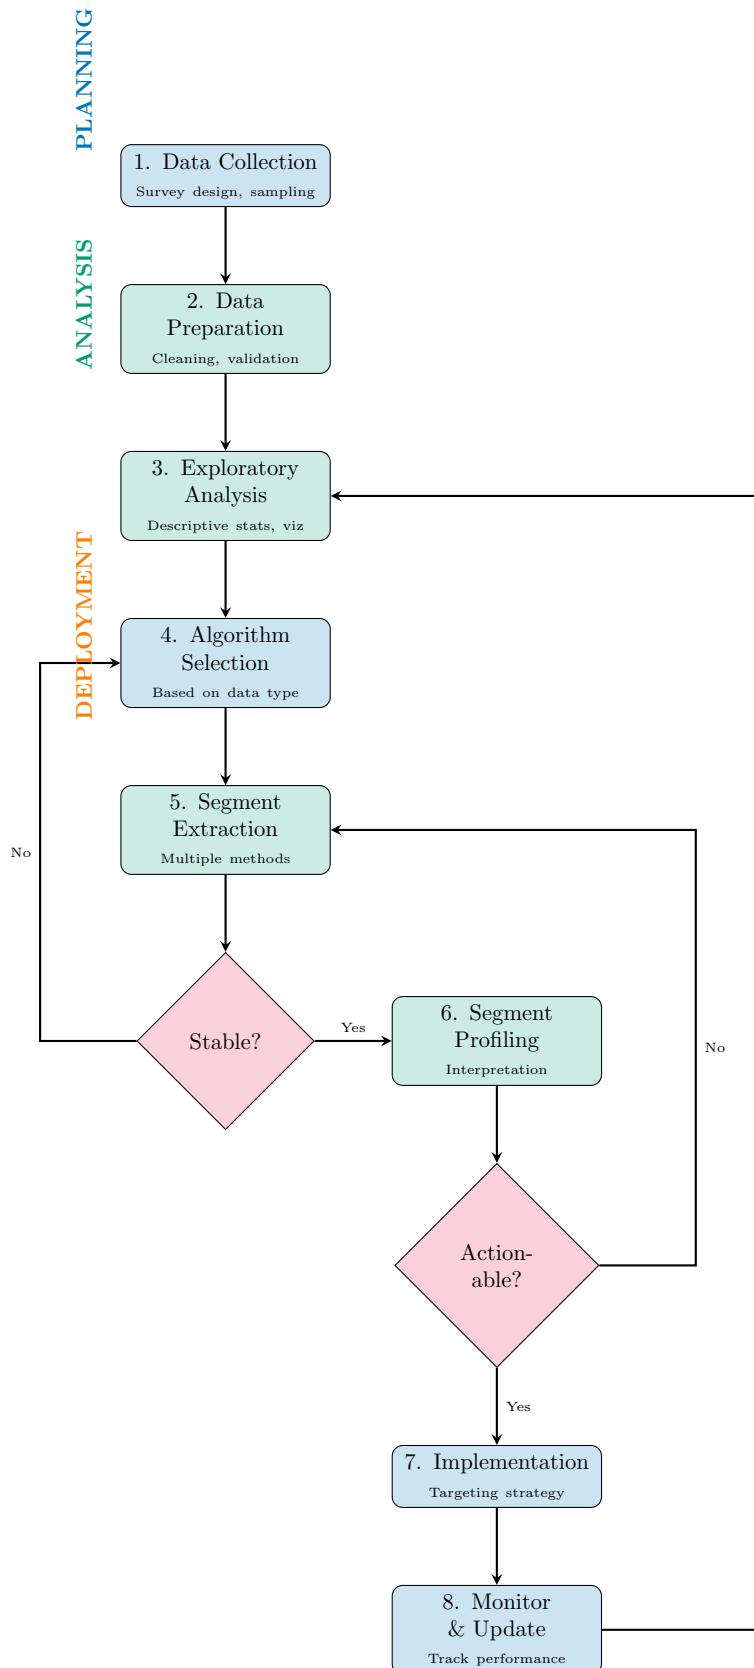


Figure 11. Complete market segmentation workflow from data collection to implementation.⁵¹
Iteration loops allow refinement when stability or actionability criteria are not met.

11.2 Decision Tree for Algorithm Selection

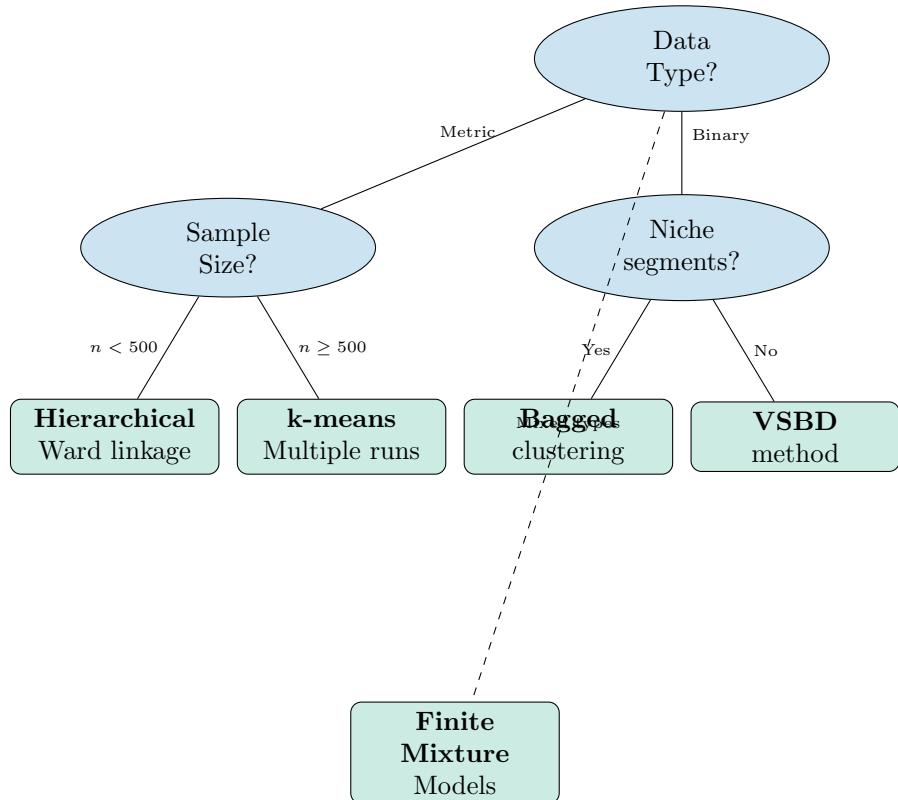


Figure 12: Decision tree for selecting appropriate segmentation algorithm based on data characteristics and business requirements.

11.3 Stage-by-Stage Checklists

Stage 1: Data Preparation Checklist

- Survey design reviewed by domain experts
- Sample size adequate ($n \geq 200$ minimum, $n \geq 500$ preferred)
- Missing data handled appropriately
- Outliers identified and investigated
- Variables standardized if on different scales
- Response style bias assessed
- Multicollinearity checked ($VIF < 10$)
- Data saved in analysis-ready format

Stage 2: Algorithm Selection Checklist

- Data type documented (binary/ordinal/metric/mixed)
- Appropriate distance measure selected
- Multiple algorithms chosen for comparison
- Algorithm assumptions match data characteristics
- Computational feasibility verified
- Software implementation tested
- Random seed set for reproducibility

Stage 3: Segment Extraction Checklist

- Multiple k values tested (typically $k = 2$ to $k = 8$)
- k-means run 50+ times with different initializations
- Dendrogram inspected for hierarchical methods
- Scree plot generated and examined
- Global stability analysis completed ($b \geq 50$ bootstrap samples)
- Segment-level stability (SLSw) computed
- Optimal k selected based on stability, not just fit
- Results documented with reproducibility details

Stage 4: Segment Validation Checklist

- Segments differ significantly on segmentation variables
- Segments interpretable with clear narratives
- Segments not artifacts of response style
- Segment sizes viable for business ($\geq 5\%$ preferred)
- Segments accessible through marketing channels
- Segments stable in SLSA analysis
- Segments validated on holdout sample if possible
- Stakeholder buy-in obtained

11.4 Common Pitfalls and Solutions

Table 5: Common Segmentation Pitfalls and Recommended Solutions

Pitfall	Consequence	Solution
Using only one algorithm	Biased solution reflecting algorithm assumptions	Test 3+ different methods; compare results
Selecting k by elbow method alone	Artificial segment count	Prioritize stability analysis over fit measures
Ignoring response style bias	Segments reflect survey artifacts, not preferences	Screen for acquiescence patterns; exclude biased segments
Accepting unstable segments	Targeting non-reproducible groups	Require median Jaccard > 0.7 in SLSw analysis
Creating too many segments	Unactionable, overlapping groups	Limit to 3-5 segments for practical targeting
Not standardizing variables	Distance dominated by high-variance features	Always standardize metric variables before clustering
Single random initialization (k-means)	Local optima, non-optimal solution	Run 50+ times; select best WCSS
Overfitting to sample	Poor generalization	Validate on holdout data; assess SLSA stability

11.5 Segment Profiling Best Practices

Creating Actionable Segment Profiles

A complete segment profile should include:

1. Demographic Description

- Age distribution, gender split, education level
- Income brackets, household composition
- Geographic concentration

2. Behavioral Characteristics

- Product/service usage patterns
- Purchase frequency and monetary value
- Channel preferences (online, retail, mobile)

3. Psychographic Insights

- Values, attitudes, lifestyle preferences
- Motivations and pain points

- Brand perceptions and preferences

4. Segment Naming

- Memorable, descriptive names (not "Segment 1")
- Captures essence of consumer group
- Examples: "Budget Seekers", "Premium Loyalists", "Convenience Focused"

5. Size and Value

- Absolute size (number of consumers)
- Percentage of total market
- Current and potential revenue
- Lifetime value estimates

11.5.1 Python Implementation: Automated Profiling Report

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 from scipy import stats
5
6 def generate_segment_profile_report(data, segment_col='segment',
7                                     segmentation_vars=None,
8                                     demographic_vars=None):
9     """
10     Generate comprehensive segment profiling report.
11     """
12     segments = data[segment_col].unique()
13     n_segments = len(segments)
14
15     report = {}
16
17     print("=" * 80)
18     print("SEGMENT PROFILING REPORT")
19     print("=" * 80)
20
21     for seg in sorted(segments):
22         seg_data = data[data[segment_col] == seg]
23         pop_size = len(seg_data)
24         pop_pct = 100 * pop_size / len(data)
25
26         print(f"\n{'='*80}")
27         print(f"SEGMENT {seg+1}")
28         print(f"{'='*80}")
29         print(f"Size: {pop_size:,} consumers ({pop_pct:.1f}% of total)")
30
31         seg_report = {
32             'size': pop_size,

```

```

33         'percentage': pop_pct
34     }
35
36     # Segmentation variables profile
37     if segmentation_vars:
38         print(f"\nSegmentation Variables:")
39         print("-" * 40)
40
41         seg_means = seg_data[segmentation_vars].mean()
42         pop_means = data[segmentation_vars].mean()
43
44         for var in segmentation_vars:
45             seg_val = seg_means[var]
46             pop_val = pop_means[var]
47             diff = seg_val - pop_val
48             pct_diff = 100 * diff / pop_val if pop_val != 0 else 0
49
50             print(f" {var:20s}: {seg_val:.2f} "
51                   f"(pop: {pop_val:.2f}, "
52                   f"{['↑' if diff > 0 else '↓']}{abs(pct_diff):5.1f}%)")
53
54     # Demographics profile
55     if demographic_vars:
56         print(f"\nDemographic Profile:")
57         print("-" * 40)
58
59         for var in demographic_vars:
60             if data[var].dtype in ['int64', 'float64']:
61                 # Numeric variable
62                 seg_val = seg_data[var].mean()
63                 pop_val = data[var].mean()
64                 print(f" {var:20s}: {seg_val:.1f} (pop: {pop_val:.1f})")
65             else:
66                 # Categorical variable
67                 seg_dist = seg_data[var].value_counts(normalize=True)
68                 print(f" {var}:")
69                 for cat, pct in seg_dist.head(3).items():
70                     print(f"    {cat}: {100*pct:.1f}%")
71
72     # Statistical distinctiveness
73     print(f"\nStatistical Tests:")
74     print("-" * 40)
75
76     if segmentation_vars:
77         # ANOVA F-statistic
78         other_data = data[data[segment_col] != seg]
79
80         for var in segmentation_vars[:5]: # Top 5 variables
81             seg_vals = seg_data[var].values
82             other_vals = other_data[var].values
83
84             f_stat, p_val = stats.f_oneway(seg_vals, other_vals)
85
86             print(f" {var:20s}: F={f_stat:.2f}, "
87                   f"p={'<0.001' if p_val < 0.001 else f'{p_val:.3f}'}")
88

```

```

89         report[f'segment_{seg+1}'] = seg_report
90
91     return report
92
93 # Example usage
94 np.random.seed(42)
95
96 # Generate sample data
97 n_total = 1000
98
99 # Create 3 segments with distinct profiles
100 seg1_data = pd.DataFrame({
101     'price_sensitivity': np.random.normal(8, 1, 350),
102     'quality_importance': np.random.normal(3, 1, 350),
103     'brand_loyalty': np.random.normal(2, 1, 350),
104     'age': np.random.normal(35, 10, 350),
105     'income': np.random.normal(45000, 10000, 350),
106     'segment': 0
107 })
108
109 seg2_data = pd.DataFrame({
110     'price_sensitivity': np.random.normal(3, 1, 400),
111     'quality_importance': np.random.normal(8, 1, 400),
112     'brand_loyalty': np.random.normal(7, 1, 400),
113     'age': np.random.normal(45, 10, 400),
114     'income': np.random.normal(75000, 15000, 400),
115     'segment': 1
116 })
117
118 seg3_data = pd.DataFrame({
119     'price_sensitivity': np.random.normal(5, 1, 250),
120     'quality_importance': np.random.normal(5, 1, 250),
121     'brand_loyalty': np.random.normal(4, 1, 250),
122     'age': np.random.normal(28, 5, 250),
123     'income': np.random.normal(55000, 12000, 250),
124     'segment': 2
125 })
126
127 data = pd.concat([seg1_data, seg2_data, seg3_data], ignore_index=True)
128
129 # Generate report
130 segmentation_vars = ['price_sensitivity', 'quality_importance', 'brand_loyalty']
131 demographic_vars = ['age', 'income']
132
133 report = generate_segment_profile_report(data,
134                                         segmentation_vars=segmentation_vars,
135                                         demographic_vars=demographic_vars)

```

11.6 From Segmentation to Targeting Strategy

Strategic Segment Selection

Not all segments deserve equal investment. Prioritize based on:

Attractiveness Criteria:

1. **Size:** Large enough to be profitable
2. **Growth:** Expanding vs. declining segments
3. **Profitability:** High lifetime value potential
4. **Accessibility:** Reachable through available channels
5. **Stability:** Reproducible across bootstrap samples
6. **Differential response:** Unique needs requiring different offerings

Table 6: Segment Prioritization Matrix Example

Segment	Size (0-10)	Value (0-10)	Stability (0-10)	Access (0-10)	Total	Priority
Premium Loyalists	6	10	9	8	33	1
Convenience Focused	8	7	8	9	32	2
Budget Seekers	9	4	6	7	26	3
Niche Enthusiasts	2	9	5	4	20	4

11.7 Reporting to Stakeholders

Effective Segmentation Presentation

Executive Summary (1 page):

- Number of segments identified
- Brief characterization of each
- Recommended primary and secondary targets
- Expected business impact

Detailed Profiles (2-3 pages per segment):

- Persona-style description with representative quote
- Visual profile (bar charts, radar plots)
- Demographics, behaviors, attitudes
- Size and value metrics
- Targeting recommendations

Technical Appendix:

- Methodology description
- Algorithms tested and selection rationale
- Stability analysis results
- Statistical validation
- Reproducibility details

11.8 Software Implementation Guide

Table 7: Python vs R for Market Segmentation

Aspect	Python	R
Hierarchical	<code>scipy.cluster.hierarchy</code>	<code>hclust()</code> , <code>agnes()</code>
k-means	<code>sklearn.cluster.KMeans</code>	<code>kmeans()</code> , <code>cclust()</code>
Finite Mixtures	<code>sklearn.mixture.GaussianMixture</code>	<code>flexmix</code> , <code>mclust</code>
SOM	<code>minisom</code> package	<code>kohonen</code> package
Stability	Custom implementation	<code>fpc::clusterboot()</code>
Visualization	<code>matplotlib</code> , <code>seaborn</code>	<code>ggplot2</code> , base plots
Ecosystem	General purpose, ML focus	Statistical focus, segmentation-specific packages

12 Final Recommendations**Key Success Factors****1. Prioritize Stability Over Fit**

- Reproducible segments more valuable than perfect fit to one sample
- Bootstrap stability analysis is non-negotiable
- Accept that consumer data lacks natural clusters

2. Use Multiple Methods

- Compare 3+ algorithms (e.g., Ward's, k-means, finite mixtures)
- Look for segments appearing across multiple solutions
- Document sensitivity to algorithm choice

3. Focus on Business Actionability

- Statistical optimality secondary to business usability

- Prefer 3-5 segments over 8+ for practical targeting
- Ensure segments inform specific marketing actions

4. Validate Thoroughly

- Holdout validation when sample size permits
- Check for response style artifacts
- Confirm segments differ on outcome variables (sales, satisfaction)

5. Monitor and Update

- Consumer preferences evolve; segments aren't permanent
- Re-run analysis annually or when market shifts
- Track segment migration over time

Questions for Continuous Improvement

After implementing segmentation strategy:

1. Are we successfully reaching each target segment through chosen channels?
2. Do segments respond differently to our marketing mix as predicted?
3. What is the actual lifetime value of each segment vs. projections?
4. Are segment sizes stable or shifting over time?
5. Should we refine segmentation variables based on learnings?
6. Do emerging consumer trends suggest need for re-segmentation?

13 Conclusion: The Art and Science of Segmentation

Market segmentation combines rigorous statistical methodology with business intuition. This comprehensive guide has demonstrated that:

1. **Algorithm choice matters:** Different methods impose different structural assumptions, shaping results
2. **Stability is paramount:** Reproducibility across bootstrap samples indicates meaningful segments
3. **No universal "best":** Context-dependent selection based on data characteristics and business objectives
4. **Validation is essential:** Multiple complementary approaches reduce risk of spurious solutions

5. **Actionability trumps sophistication:** Simpler, interpretable segments often outperform complex models in practice

The methods and workflows presented equip practitioners to conduct rigorous, reproducible market segmentation analysis that translates statistical findings into competitive business advantage.

*The goal is not to find perfect segments,
but to identify reproducible, actionable consumer groups
that enable differentiated marketing strategy.*

References

- [1] Dolnicar, S., Grün, B., and Leisch, F. (2018). *Market Segmentation Analysis: Understanding It, Doing It, and Making It Useful*. Springer.
- [2] Hennig, C. (2007). Cluster-wise assessment of cluster stability. *Computational Statistics & Data Analysis*, 52(1), 258-271.
- [3] Leisch, F. (1999). Bagged clustering. Working Paper 51, SFB Adaptive Information Systems and Modeling in Economics and Management Science, Vienna University of Economics.
- [4] Wedel, M., and Kamakura, W.A. (2000). *Market Segmentation: Conceptual and Methodological Foundations*. 2nd ed., Kluwer Academic Publishers.
- [5] Fraley, C., and Raftery, A.E. (2002). Model-based clustering, discriminant analysis, and density estimation. *Journal of the American Statistical Association*, 97(458), 611-631.
- [6] Kohonen, T. (2001). *Self-Organizing Maps*. 3rd ed., Springer Series in Information Sciences.
- [7] Ward, J.H. (1963). Hierarchical grouping to optimize an objective function. *Journal of the American Statistical Association*, 58(301), 236-244.
- [8] Gower, J.C. (1971). A general coefficient of similarity and some of its properties. *Biometrics*, 27(4), 857-871.

